


# Parallel and Distributed Processing

Enterprise Architectures for Big Data

# Single Node

Computer Node



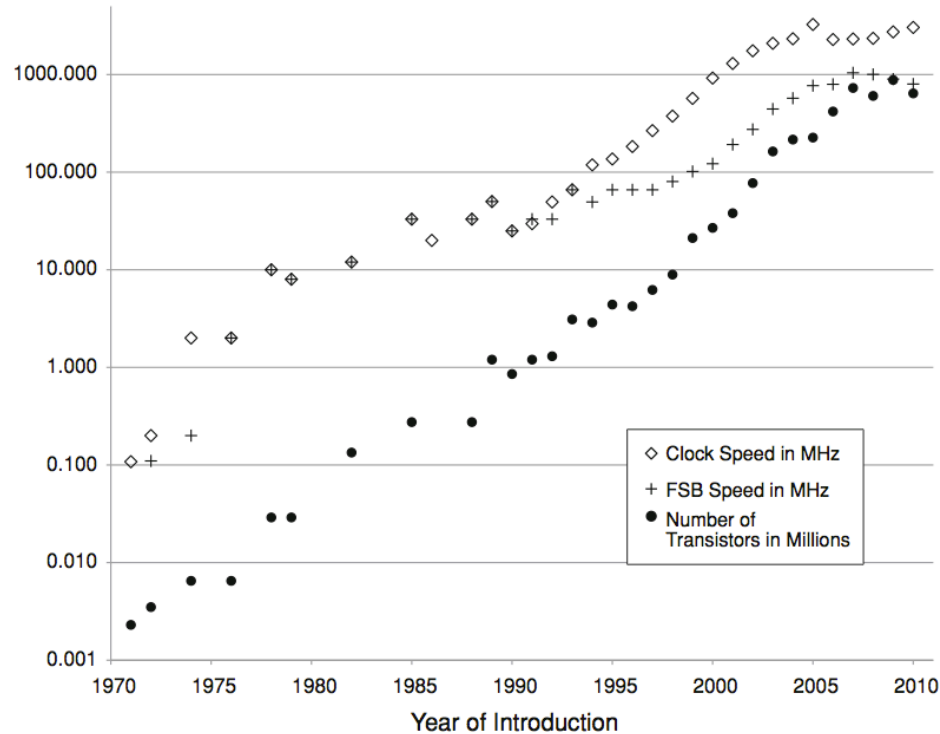
# Questions

- What is the difference between parallel computing and distributed computing?
  - Parallel Computing
    - more than one Core, CPU, GPU or Node is working on an algorithm at the same time
  - Distributed Computing
    - More than one Node is working on an algorithm at the same time
- What is the opposite of parallel computing?
  - Sequential Computing

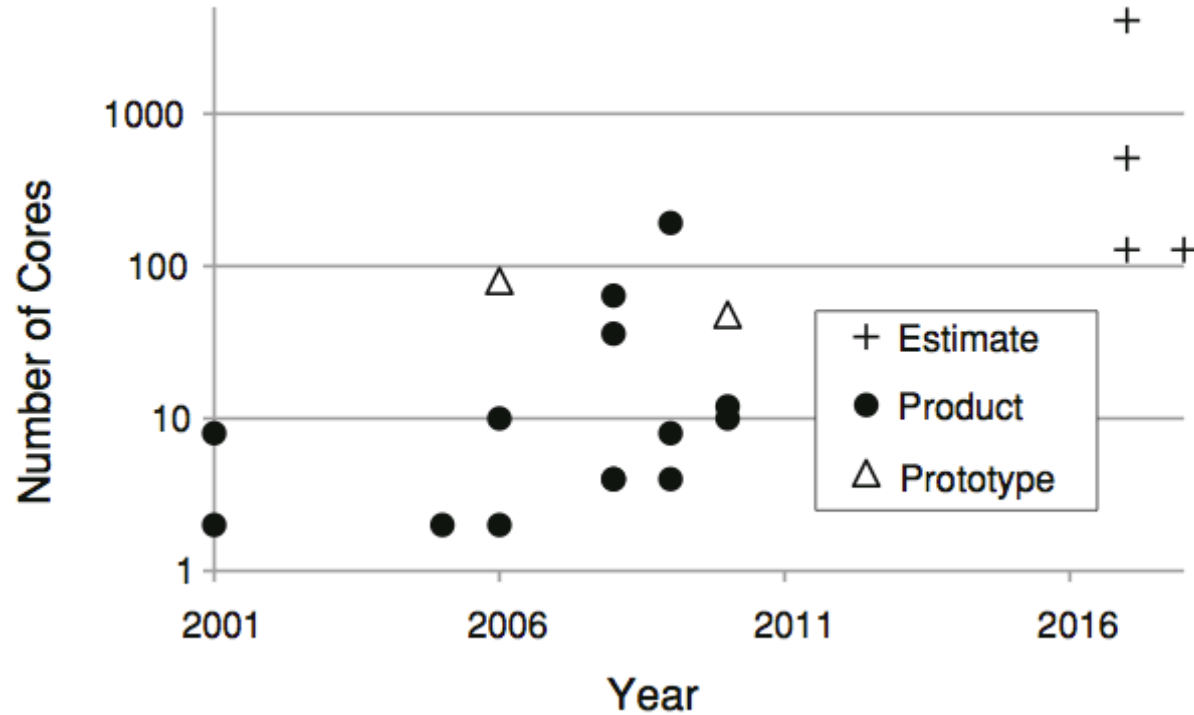
# Questions

- Why distribute computation?
  - Speedup algorithms through parallel computing
  - Fault tolerance
  - Reduce latency through nearer nodes
- Why distribute data?
  - Storage limits of single node
  - Store data near computation
  - Fault tolerance through replication
  - Reduce latency through nearer nodes
  - Improves reading speed through parallel reading


# Clock speed, Front Side Bus (FSB) speed, and number of transistors



# Development of number of cores



# Scaling Architectures: Scale up vs. Scale out



E.g.  
SAN (Storage Area Network) or  
NAS (Network Attached Storage)

# Latency Comparison Numbers

L1 cache reference	0.5ns			
L2 cache reference	7 ns			14x L1 cache
Main memory reference	100 ns			14x L2 cache, 200x L1 cache
Send 1KB over 1 Gbps network	10,000 ns	0.01 ms		
SSD seek	100,000 ns	0.1 ms		1000x memory reference
Read 4KB randomly from SSD*	150,000 ns	0.15 ms		
Read 1 MB sequentially from memory	250,000 ns	0.25 ms		
Round trip within same datacenter	500,000 ns	0.5 ms		5x SSD seek, 5,000x memory reference
Read 1 MB sequentially from SSD*	1,000,000 ns	1 ms		4X memory sequential read
Disk seek	10,000,000 ns	10 ms		20x datacenter roundtrip, 100x SSD seek, 100,000x memory reference
Read 1 MB sequentially from disk	20,000,000 ns	20 ms		80x memory, 20X SSD
Send packet CA->Netherlands->CA	150,000,000 ns	150 ms		



# Latency Comparison Numbers

## Lets multiply all these durations by a billion ( $10^9$ )

### ### Seconds:

L1 cache reference 0.5 s

One heart beat (0.5 s)

L2 cache reference 7 s

Long yawn

### ### Minutes:

Main memory reference 100 s

Brushing your teeth

### ### Hours:

Send 2K bytes over 1 Gbps network 5.5 hr

From lunch to end of work day

### ### Days

SSD seek 1.7 days

A normal weekend

Read 1 MB sequentially from memory 2.9 days

A long weekend

Round trip within same datacenter 5.8 days

A medium vacation

Read 1 MB sequentially from SSD 11.6 days

Waiting for almost 2 weeks for a delivery

### ### Months

Disk seek 16.5 weeks

A semester in university

Read 1 MB sequentially from disk 7.8 months

Almost producing a new human being

The above 2 together = 1 year

### ### Years

Send packet CA→Netherlands→CA 4.8 years

Average time it takes to complete a bachelor's degree

# The Problem of Shared State

```
x = 10
```

```
x = function1(x)
```

```
x = function2(x)
```

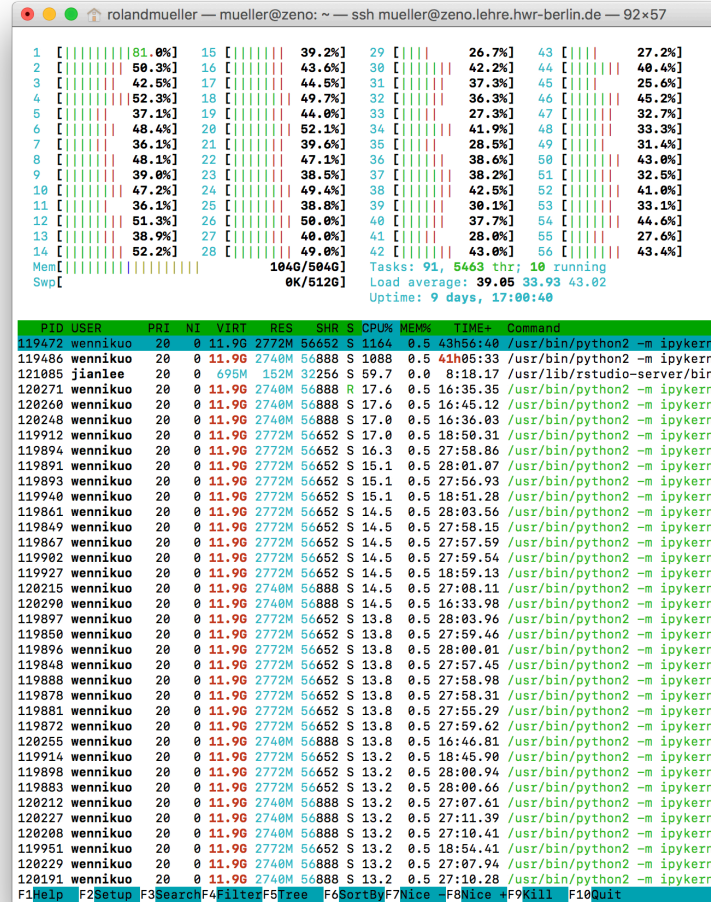
```
x = x + 10
```

- Value of x is time-dependent
- Data dependency
- Single process (sequential computing) → No problem
- Multiple parallel processes
  - Locking of value
  - Processes have to wait for each other


# Multi-Core



# htop



# Deadlock



# Pure Functions and Parallel Processing

```
x = 10
```

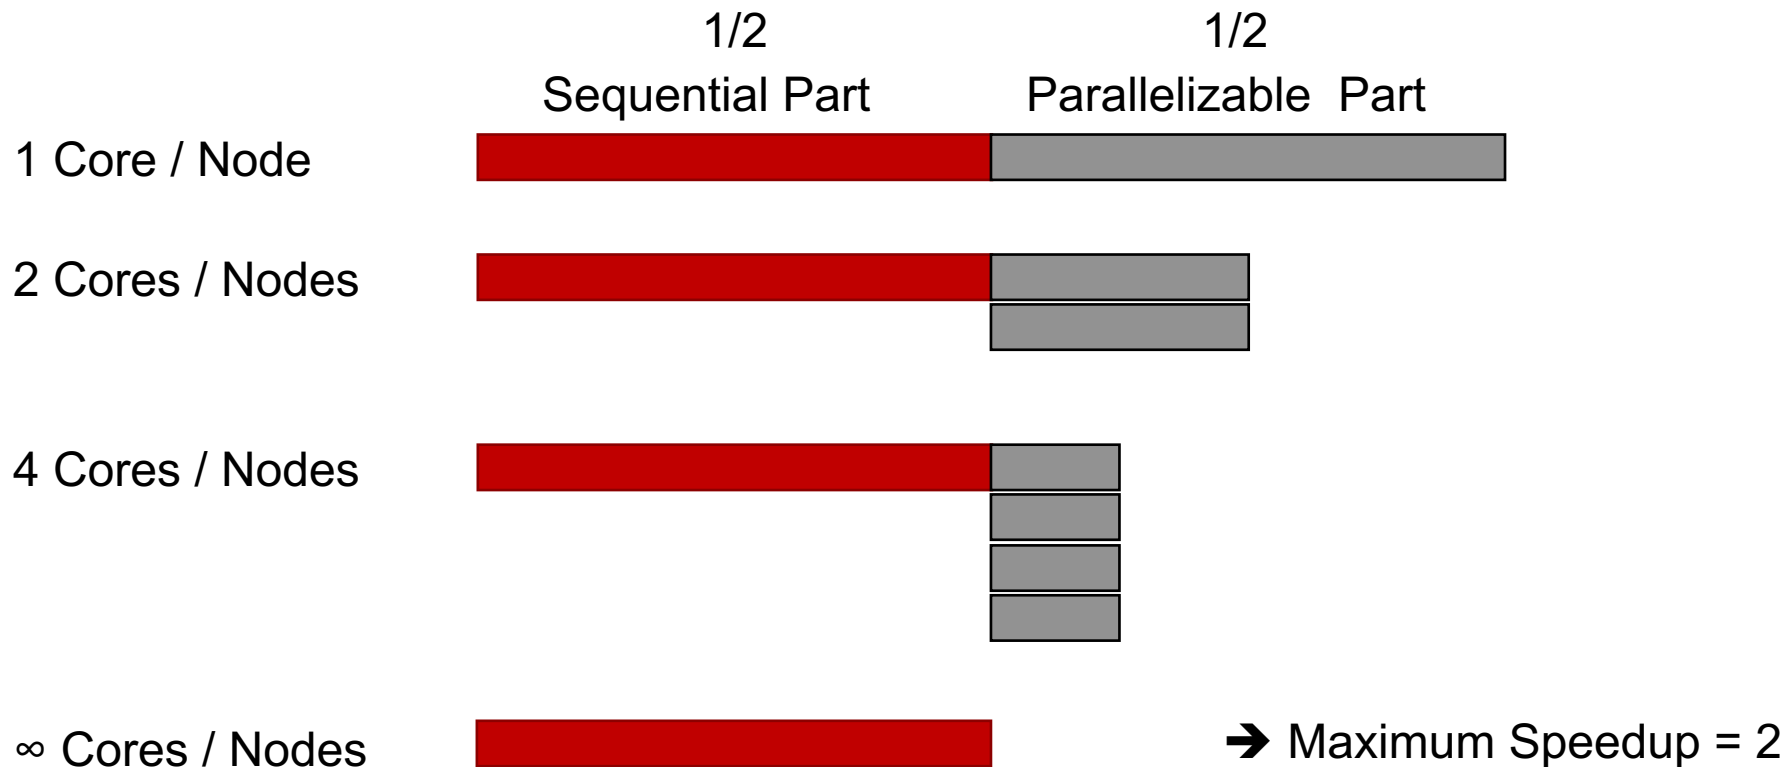
```
y = function1(x)
```

```
z = function2(x)
```

```
a = x + y
```

- Step 2 and 3 can be only parallelized if `function1` and `function2` are pure functions
  - No side effects
  - No internal state
  - Result is only determined by the input parameter

# Amdahl's Law Visually Explained



# Ahmdal's Law

- $T_1$  Total runtime of a Task (with one CPU)
- $T_s$  Runtime of sequential part (can not be parallelized)
- $T_p$  Runtime of parallelizable part

$$T_1 = T_s + T_p$$

- $f$  Parallelizable fraction of program

$$f = \frac{T_p}{T_1} \quad \text{with } 0 \leq f \leq 1$$

- $n$  CPUs / Cores / Nodes

- $T_n(n)$  Runtime with  $n$  CPUs / Cores / Nodes

$$T_n(n) = T_s + \frac{T_p}{n}$$

$$T_n(n) = (1 - f)T_1 + \frac{f}{n}T_1$$

- $S(n)$  Speedup with  $n$  CPUs

$$S(n) = \frac{T_1}{T_n(n)}$$

$$S(n) = \frac{T_1}{(1-f)T_1 + \frac{f}{n}T_1} = \frac{T_1}{T_1((1-f) + \frac{f}{n})}$$

$$S(n) = \frac{1}{(1-f) + \frac{f}{n}}$$



# Ahmdal's Law

- Speedup  $S(n)$  with  $n$  processors

$$S(n) = \frac{1}{(1-f) + \frac{f}{n}}$$


- Ideal Speedup:  $f=1$        $\rightarrow S(n)=n$
- Usual  $f<1$        $\rightarrow S(n)$  is bound by constant

$$\lim_{n \rightarrow \infty} S(n) = \frac{1}{1-f}$$

- For  $f=0.9$  and  $n=10$        $\rightarrow S(n)=5.3$
- For  $f=0.9$  and  $n \rightarrow \infty$        $\rightarrow S(n) = 10$

# Parallel Speedup

- Implications of Amdahl's Law:
- If the parallel portion of the algorithm is 95% ( $f=0.95$ ) the maximum speedup is 20



# Embarrassingly Parallel Problems

- Little or no dependency and no need for a lot of communication between the parallel tasks
- Amdahl's Law:  $f \approx 1$
- Examples:
  - Rendering Computer Graphics
  - K-fold Cross-Validation in Data Mining
  - Testing a Data Mining Algorithm with different parameters
  - Serving static files on a webserver to multiple users at once
  - Computer simulations comparing many independent scenarios, such as climate models
  - Tree growth step of the random forest
  - ...