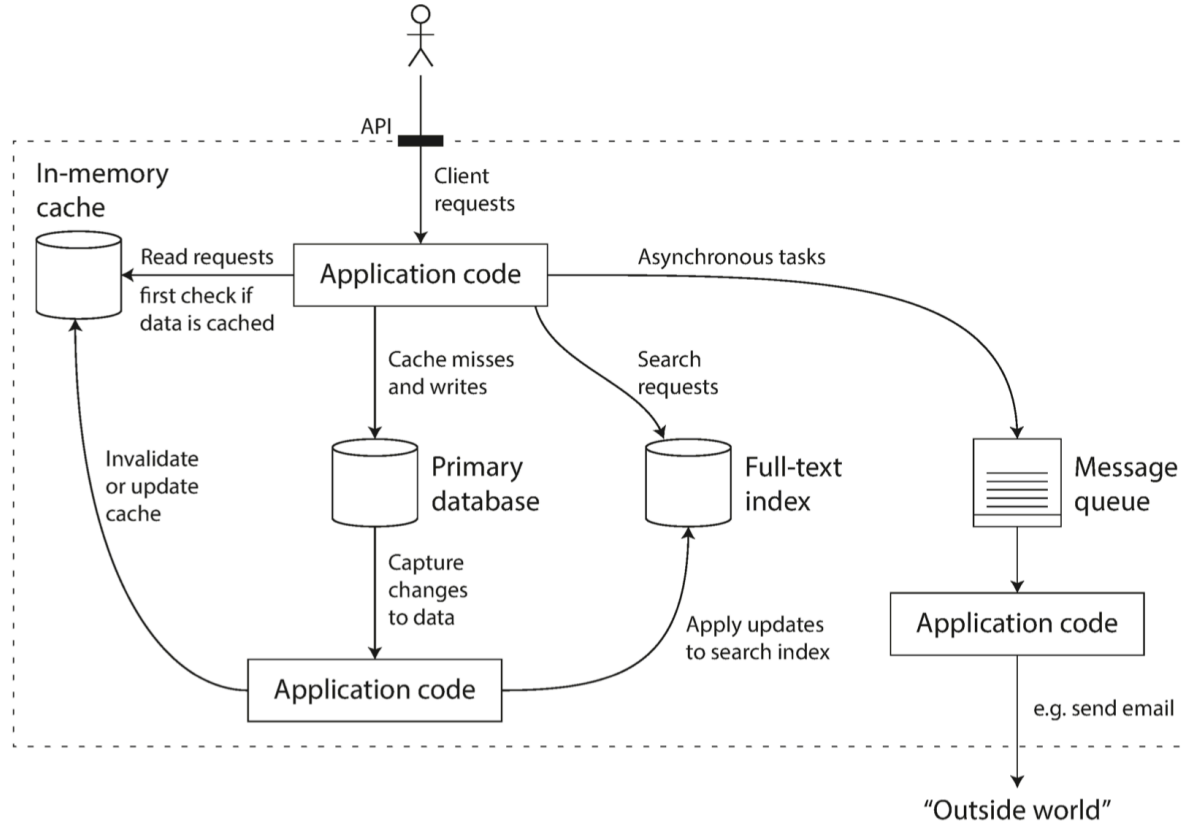


Designing Data-Intensive Applications

Enterprise Architectures for Big Data

Parts of a Modern Application



Parts of a Modern Application

■ Database

- Store data so that this or another application can find it again later

■ Cache

- Remember the result of an expensive operation, to speed up reads

■ Search index

- Allow users to search data by keyword or filter it in various ways

■ Stream processing

- Message Queue
- Send a message to another process, to be handled asynchronously

■ Batch processing

- Periodically crunch a large amount of accumulated data

Requirements

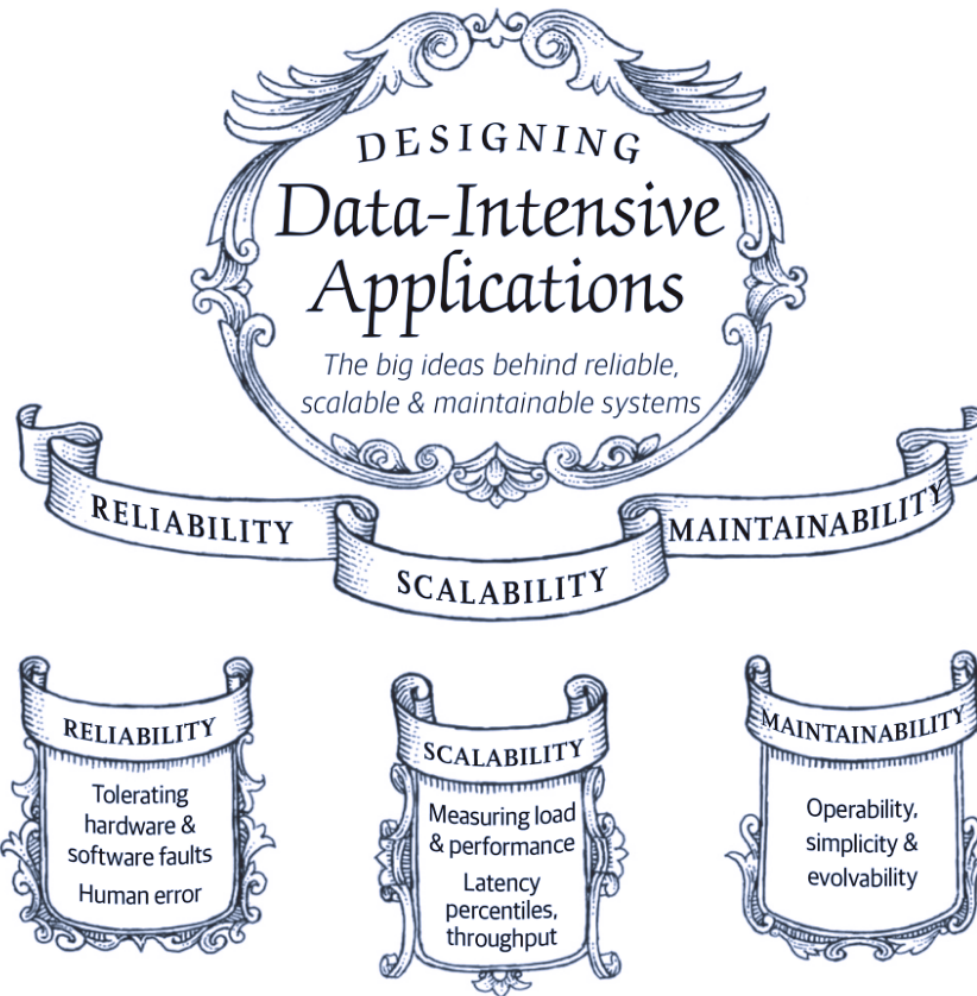
■ Functional requirements

What the system should do, such as allowing data to be stored, retrieved, searched, and processed in various ways.

■ Nonfunctional requirements

Are general properties like

- security
- reliability
- compliance
- scalability
- compatibility
- maintainability



3 main concerns of most software systems

1. Reliability

- The system should **continue to work correctly** (performing the correct function at the desired level of performance) even in the face of adversity (hardware or software faults, and even human error).

2. Scalability

- As the system grows (in data volume, traffic volume, or complexity), there should be **reasonable ways of dealing with that growth**.

3. Maintainability

- Over time, many different people will work on the system (engineering and operations, both maintaining current behavior and adapting the system to new use cases), and they should all be **able to work on it productively**.

Reliability

Reliability

- The application executes the functions that the user expected.
- It can tolerate that the users are making mistakes or using the software in unexpected ways.
- Its performance is good enough for the required use case, under the expected load and data volume.
- The system prevents any unauthorized access and abuse.

- How important is reliability?
- How can a system fail?

Fault Types

- Hardware Faults
 - E.g. Hard disks:
Mean time to failure (MTTF): 10-50 years
→ Cluster of 10,000 disks → 1 disk dies per day
- Software Errors
- Human Errors

- How are hardware and software errors different?
- How to make a system more fault tolerant?

Fault Tolerant Systems

- Strategies of fault tolerant systems
 - Redundancy
 - Testing
 - Well designed abstraction (API)
 - Isolation
 - Decoupling
 - Measuring
 - Monitoring
 - Quick recovery
 - Backup
 - Good management practice and training

Scalability

Scalability

- Definition: the system's ability to cope with increased load
- Not a one-dimensional label
- Meaningless to say “X is scalable” or “Y doesn't scale”

Rather say:

- “If the system grows in a particular way, what are our options for coping with the growth?”
- “How can we add computing resources to handle the additional load?”

Describing Performance

Two different ways to conceptualize performance:

1. When you **increase** a **load parameter** and keep the system **resources** (CPU, memory, network bandwidth, etc.) **unchanged**, how is the **performance** of your system affected?
2. When you **increase** a **load parameter**, how much do you need to **increase** the **resources** if you want to keep **performance unchanged**?

Load Parameters

- Load parameters describe the load
- Depends on your application
- KPIs
 - requests per second to a web server
 - ratio of reads to writes in a database
 - number of simultaneously active users in a chat room

Performance Measurement

Throughput


- Definition:
 - Number of records we can process per second,
 - Alternatively: the total time it takes to run a job on a dataset of a certain size
- Important for Batch Systems (like Hadoop)

Performance Measurement

Response time

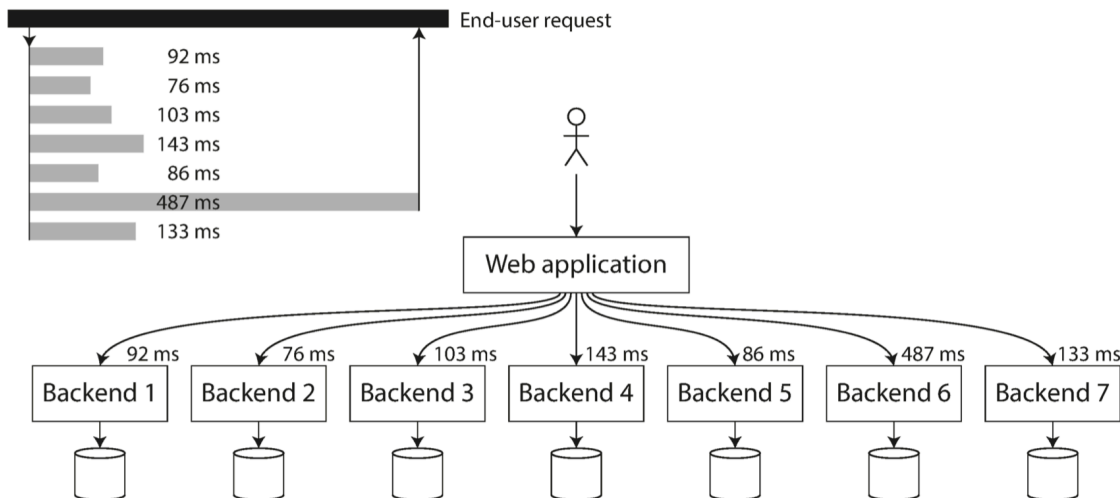
- Definition:
 - Time between a client sending a request and receiving a response
- Important for Online Systems (like an E-Shop)
- Latency vs. response time:
 - Response time: what the client sees
 - Besides the actual time to process the request (the service time), it includes network delays and queueing delays.
- Effect of Response time:
 - Amazon: 100 ms increase in response time reduces sales by 1%
 - 1 s slowdown reduces a customer satisfaction metric by 16%

Response time



- Aggregations:
 - Not just the average
 - Percentile
- **Tail latencies**, e.g. 99.9th percentile
- Customers with the slowest requests → most data on their accounts → made many purchases → most valuable customers


Response time



■ Tail latency amplification

- It only takes a small number of slow requests to hold up the processing of subsequent requests
- Queueing delays
- Head-of-line blocking

Scaling Architectures: Scale up vs. Scale out



E.g.
SAN (Storage Area Network) or
NAS (Network Attached Storage)

Distributed Systems

- Stateless vs. stateful data service
 - Stateless services: Distributing across multiple machines fairly straightforward
 - Stateful services: Distributed setup can introduce a lot of complexity
- Tradeoff: Scalability vs. Simplicity (Maintainability)
- Therefore, common wisdom until recently was to **keep your database on a single node** (scale up) until scaling cost or high-availability requirements forced you to make it distributed.
- New tools and abstractions → Change of the tradeoffs

No generic, one-size-fits-all scalable architecture

- The scaling problem may be
 - the volume of reads,
 - the volume of writes,
 - the volume of data to store,
 - the complexity of the data,
 - the response time requirements,
 - the access patterns,
 - or some mixture of all of these plus many more issues.
- An architecture that scales well for a particular application is built around **assumptions**
 - which operations will be frequent, and which will be rare → **load parameters**
- Architectures of systems that operate at large scale are usually **highly specific to the assumptions**
 - E.g. in an early-stage **startup** it's usually more important to be able to iterate quickly on product features than it is to scale to some hypothetical future load.
- **No generic, one-size-fits-all scalable architecture**
 - No Magic Scaling Sauce
 - However, scalable architecture are based on general-purpose building blocks and arranged in familiar patterns

Maintainability

Maintainability

■ Operability

- Make it easy for operations teams to keep the system running smoothly.

■ Simplicity

- Make it easy for new engineers to understand the system, by removing as much complexity as possible from the system.
- Not the same as simplicity of the user interface

■ Evolvability

- Make it easy for engineers to make changes to the system in the future, adapting it for unanticipated use cases as requirements change.

Operability

Keeping a software system running smoothly

- **Monitoring:** Providing visibility into the runtime behavior and internals of the system
- **Automation** (Continuous Integration (CI), Continuous Delivery (CD))
- **Integration** with standard tools
- **Avoiding dependency** on individual machines
- Good **documentation**
- Easy-to-understand **operational model**
- Providing **good default behavior**, optionally override defaults when needed
- **Self-healing** where appropriate, but also giving administrators manual control over the system state when needed
- **Predictable behavior**, minimizing surprises

Simplicity

- Accidental complexity:
 - Not inherent in the problem but arises only from the implementation
- Best tools for removing accidental complexity: **Abstraction**
- Good Abstraction:
 - can hide a great deal of implementation detail behind a clean, simple-to-understand façade.
 - can also be used for a wide range of different applications.
 - Reuse is more efficient than reimplementing a similar thing multiple times
 - Leads to higher-quality software, as quality improvements in the abstracted component benefit all applications that use it.
 - Finding good abstractions is hard

Evolvability

Making Change Easy

- Changing requirements
- Demand of Agility
- Influenced by:
 - Analyzability
 - Extensibility
 - Integrity
 - Portability
 - Testability
 - Changeability
- Linked to Simplicity and Abstraction
- Simple and easy-to-understand systems are easier to modify than complex ones