

# NoSQL

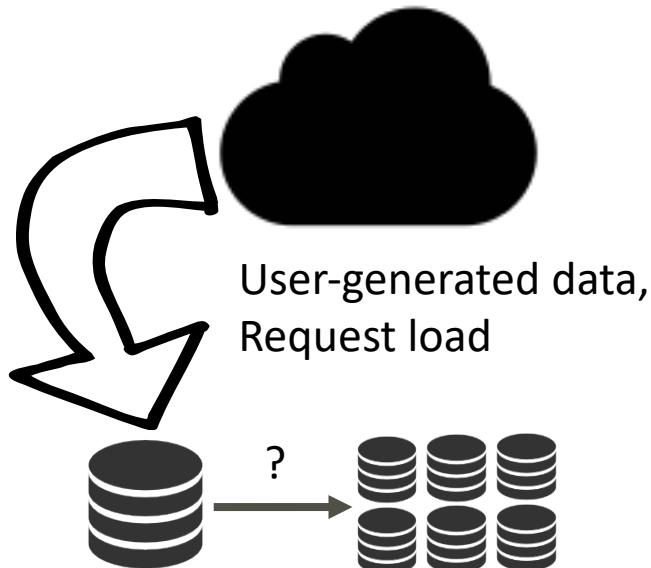
Enterprise Architectures for Big Data

# NoSQL

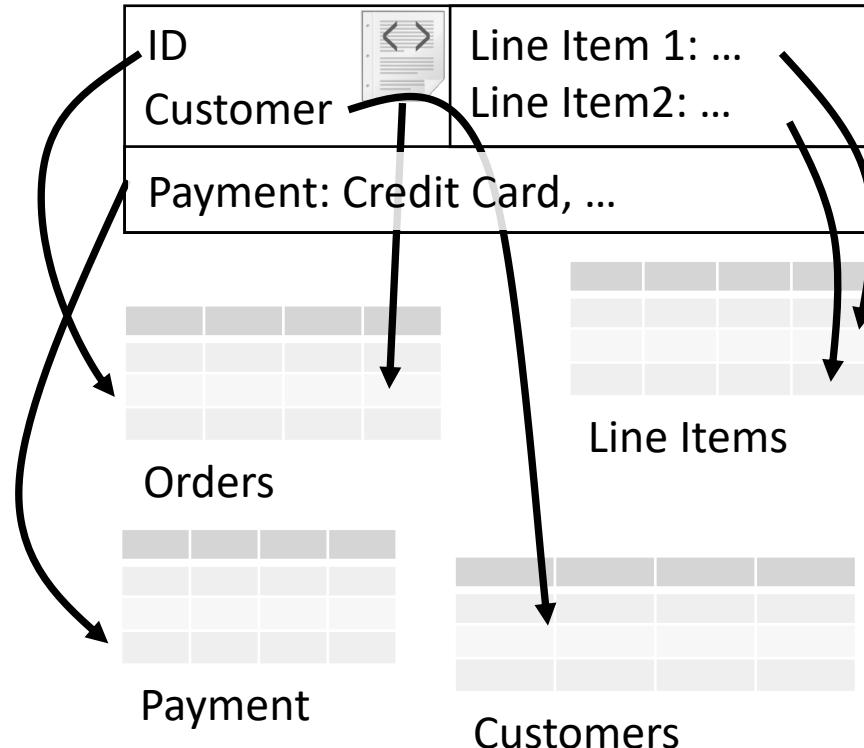
- „NoSQL“ term coined in 2009
- Interpretation: „Not Only SQL“
- Typical properties:
  - Non-relational
  - Often Open-Source
  - Schema-less (schema-free)
  - Optimized for distribution (clusters)
  - Tunable consistency

# Two main motivations

## Scalability

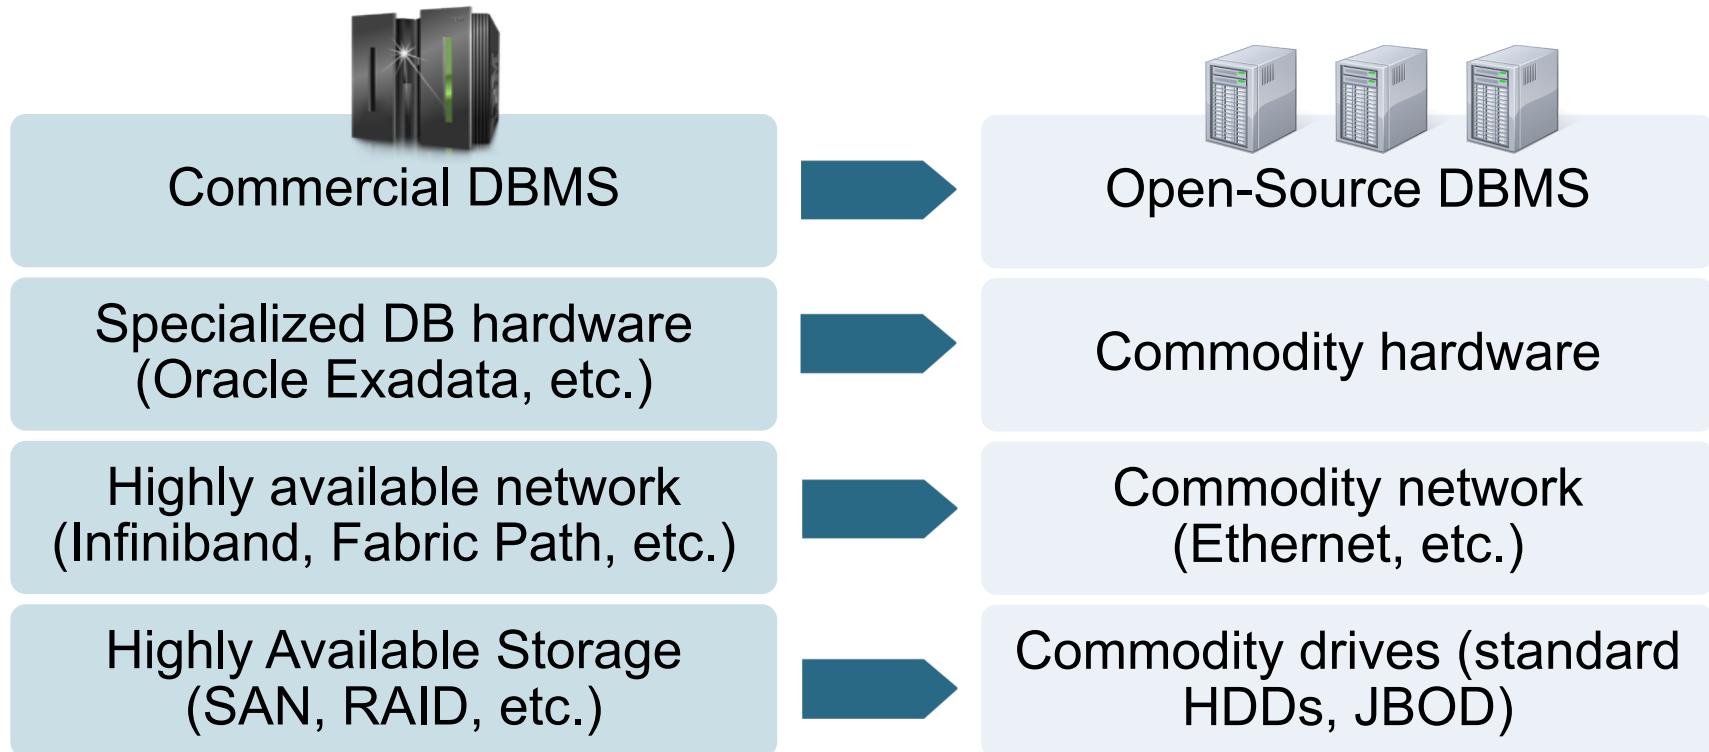


## Impedance Mismatch



# NoSQL Paradigm Shift

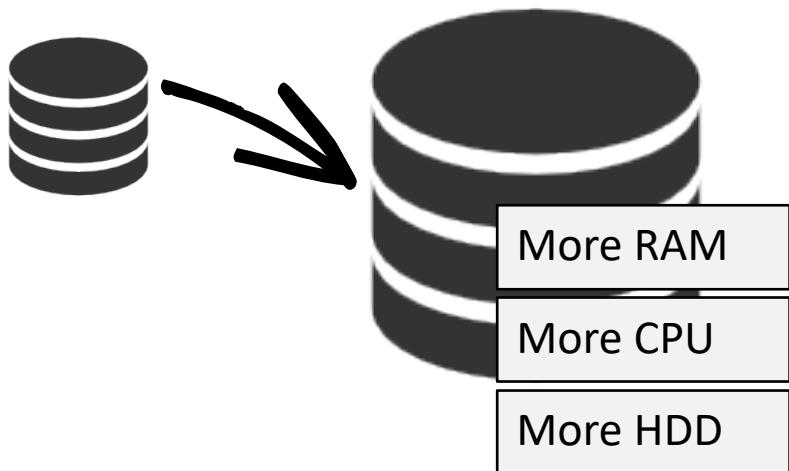
## Open Source & Commodity Hardware



# Scaling Strategies

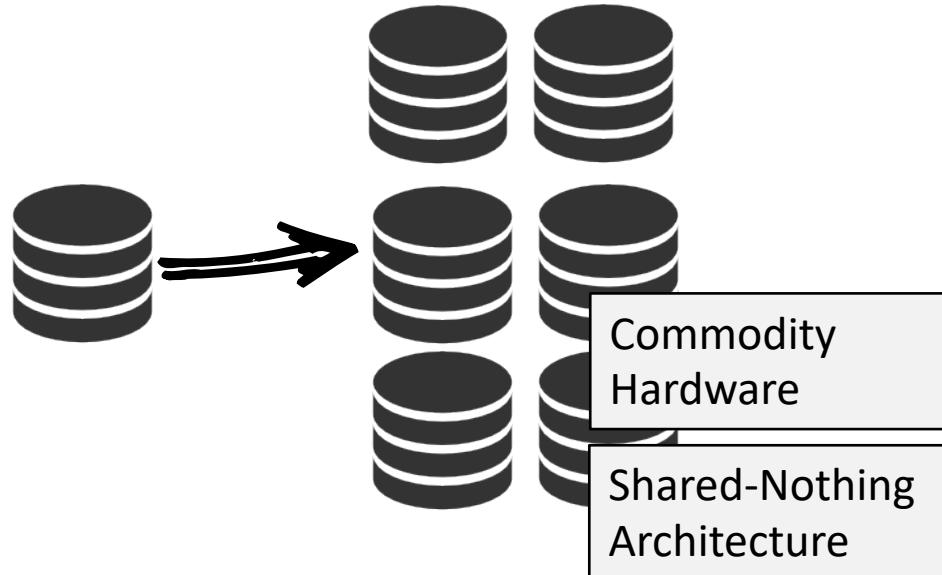
## Scale-up

- Vertical Scaling



## Scale-out

- Horizontal scaling



## Data Analytics

Analytics

Reporting

Data Mining

Data Warehouse

Big Data Platforms



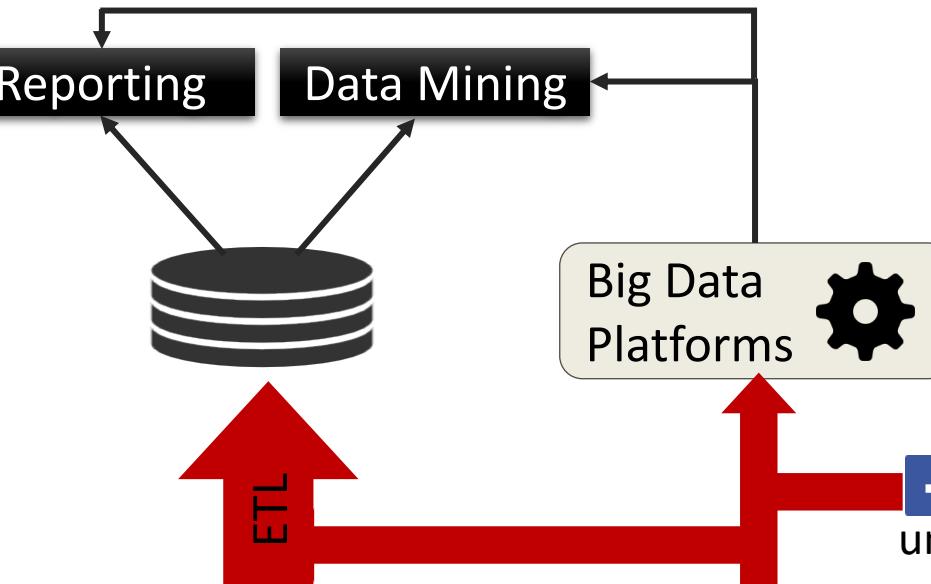
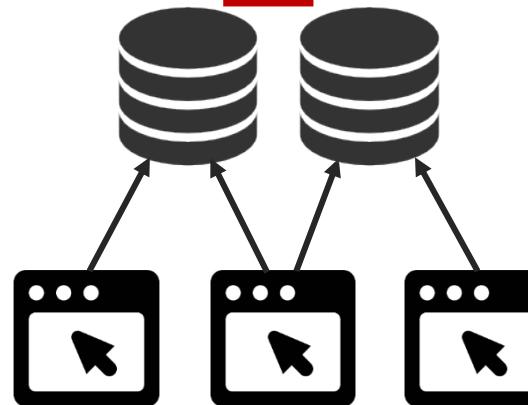
## Data Management

Operational Databases

NoSQL DBs



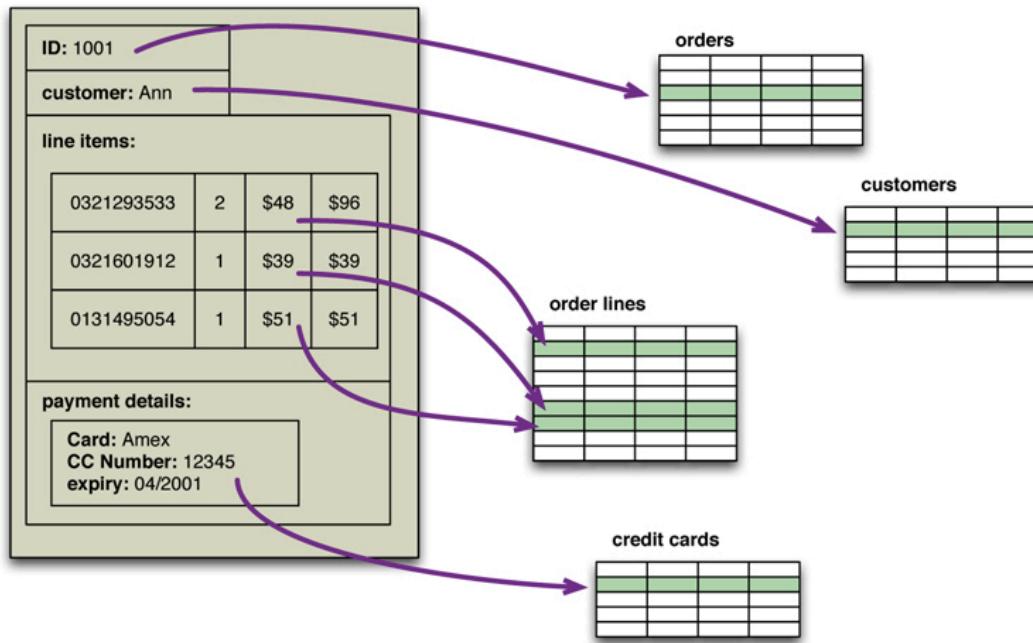
Application



unstructured

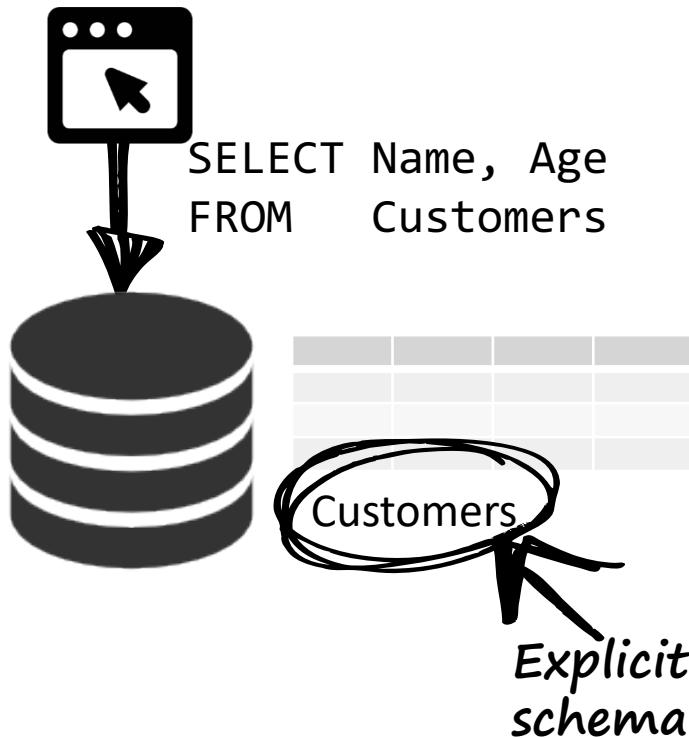
# Impedance Mismatch

- An order spans many tables in a relational database

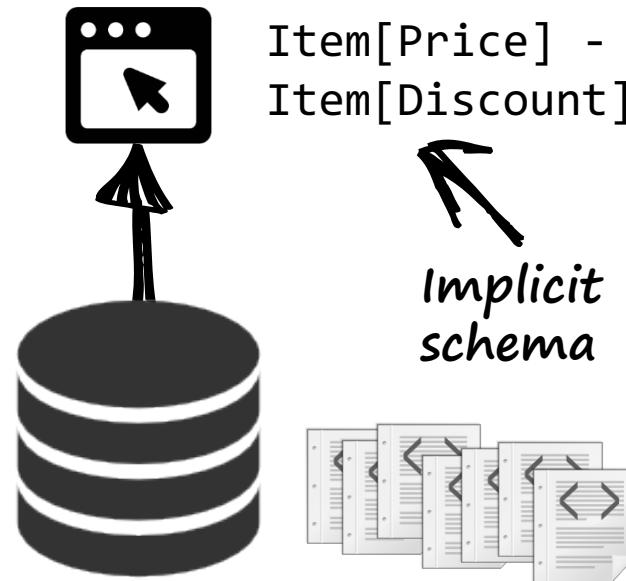


# Schema Free Data Modeling

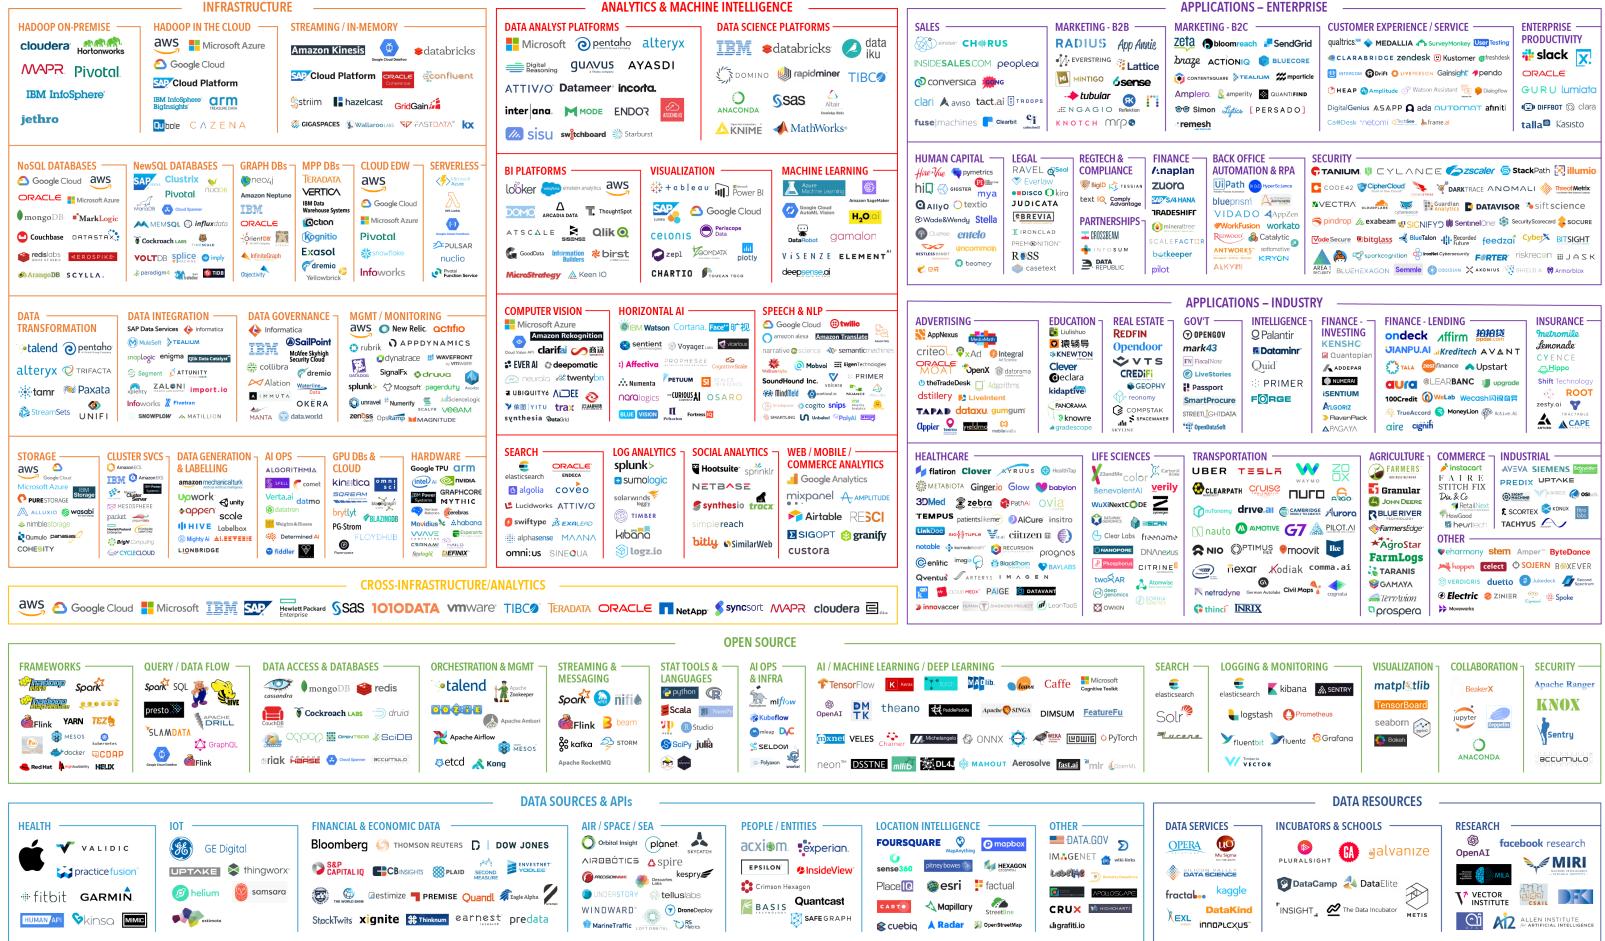
## RDBMS



## NoSQL DB



DATA & AI LANDSCAPE 2019



## INFRASTRUCTURE

### HADOOP ON-PREMISE

cloudera Hortonworks

MAPR Pivotal

IBM InfoSphere

jethro

### HADOOP IN THE CLOUD

aws Microsoft Azure

Google Cloud

SAP Cloud Platform

IBM InfoSphere BigInsights

arm TREASURE DATA

Dubole

CAZENA

### STREAMING / IN-MEMORY

Amazon Kinesis Google Cloud Dataflow databricks

SAP Cloud Platform ORACLE Coherence confluent

striim hazelcast GridGain

GIGASPCES WallarooLabs FASTDATA ix

### NoSQL DATABASES

Google Cloud aws

ORACLE Microsoft Azure

mongoDB MarkLogic

Couchbase DATASTAX

redislabs AEROSPIKE

ArangoDB SCYLLA

### NewSQL DATABASES

SAP HANA Clustrix nuoDB

MariaDB Cloud Spanner

MEMSQL influxdata

Cockroach LABS TIMESCALE

VOLTDB splice imply

paradigm4 Trafodion TIBB

### GRAPH DBs

neo4j Amazon Neptune

IBM Cloud Spanner

ORACLE OrientDB

InfiniteGraph

dremio Yellowbrick

### MPP DBs

TERADATA VERTICA

IBM Data Warehouse Systems

action

Kognitio

Exasol

dremio Yellowbrick

### CLOUD EDW

aws Microsoft Azure

Google Cloud

Microsoft Azure

Pivotal

snowflake

Infoworks

### SERVERLESS

Microsoft Azure AWS

AWS Lambda

Google Cloud Functions

PULSTAR

nuclo

Pivotal Function Service

### DATA TRANSFORMATION

talend pentaho A Health Group Company

alteryx TRIFACTA

tamr Paxata

StreamSets UNIFI

### DATA INTEGRATION

SAP Data Services Informatica

MuleSoft TEALIUM

snapLogic enigma Qlik Data Catalyst

Segment ATTUNITY

xplenty ZALONI FISIATRON

Infoworks import.io

SNOWPLOW MATILLION

### DATA GOVERNANCE

Informatica SailPoint

IBM McAfee Skyhigh Security Cloud

collibra dremio

Alation Waterline Data

IMMUTA OKERA

MANTA data.world

### MGMT / MONITORING

aws New Relic actifio

rubrik dynatrace

APPDYNAMICS WAVEFRONT by VMware

DATADOG SignalFx druva

splunk Moogsoft pagerduty

unravel Numerify scalyr

zenoss OpsRamp VEEAM

### COMPUTER VISION

Microsoft Azure

Amazon Rekognition

Cloud Vision API

clarifai 商场

sentient technologies Voyager Labs

PROPHESEE

AFFECTIVA twentybn

### HORIZONTAL AI

IBM Watson Cortana Face++ 眇视

amazon alexa

narrative science

Wolfram Alpha

Mobileye

Gridspace

cogito

## ANALYTICS & MACHINE INTELLIGENCE

### DATA ANALYST PLATFORMS

Microsoft

Pentaho A Health Group Company

Alteryx

Digital Reasoning

Guavus a Thales company

Ayasdi

Attivio Datameer

Incorta

Interana

Mode

ENDOR

Sisu

switchboard

Starburst

### DATA SCIENCE PLATFORMS

IBM databrick

DOMINO rapidm

ANACONDA

Sas

KNIME

Power BI

### BI PLATFORMS

looker

salesforce einstein analytics

aws

DOMO ARCADIA DATA

ThoughtSpot

AT SCALE SENSE

Qlik

GoodData Information Builders

birst

MicroStrategy Keen IO

Tableau

Microsoft Power BI

SAP Lumira

Google Cloud

Periscope Data

Celonis

Zepel

QOMDATA AlgoReplay

Plottly

CHARTIO Toucan Toco

### VISION & NLP

Google Cloud

amazon alexa

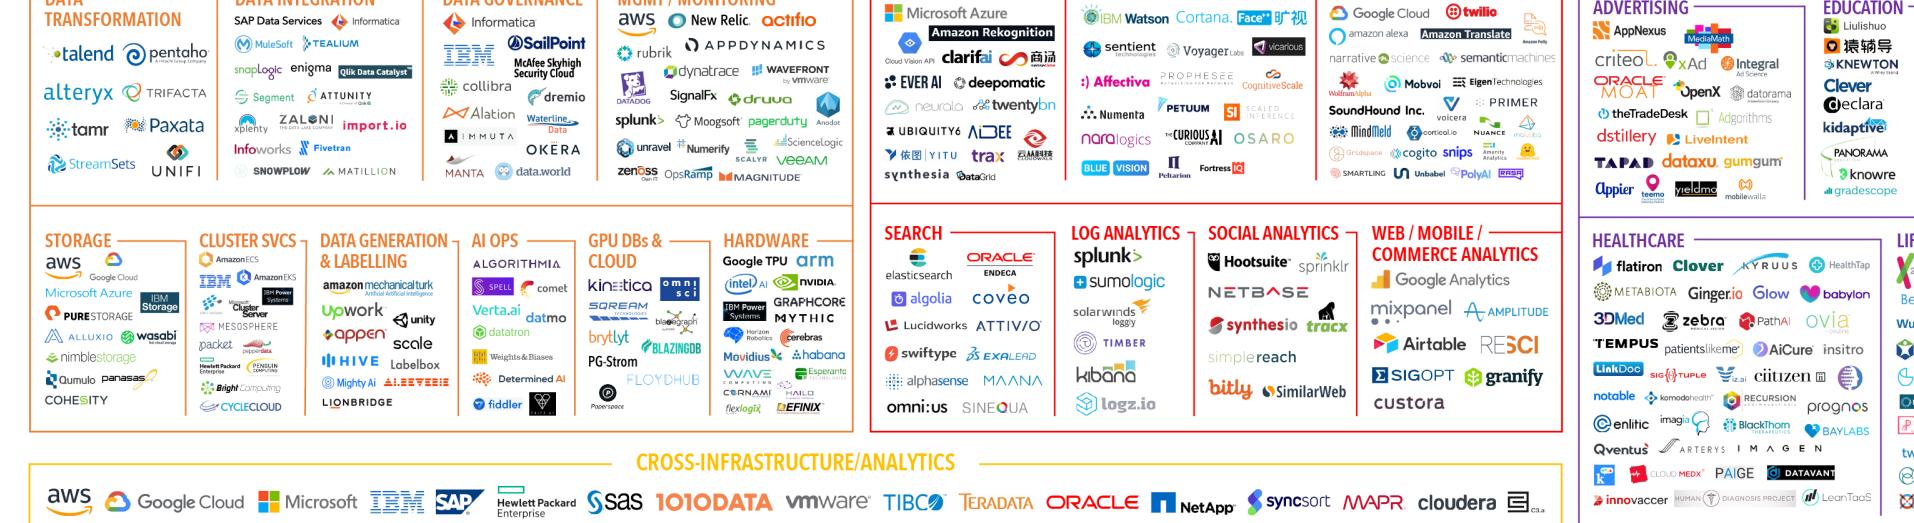
narrative science

Wolfram Alpha

Mobileye

Gridspace

cogito

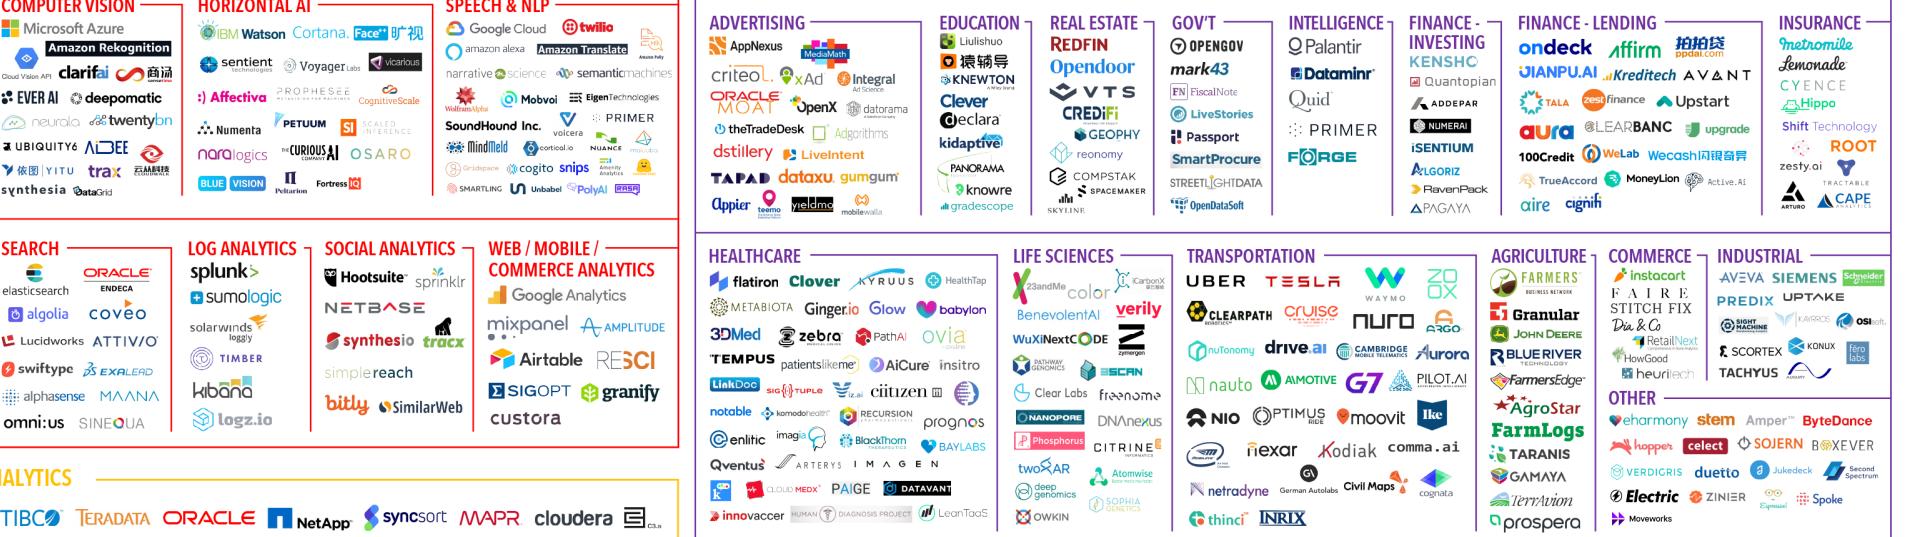


## CROSS-INFRASTRUCTURE/ANALYTICS



## DATA SOURCES & APIs





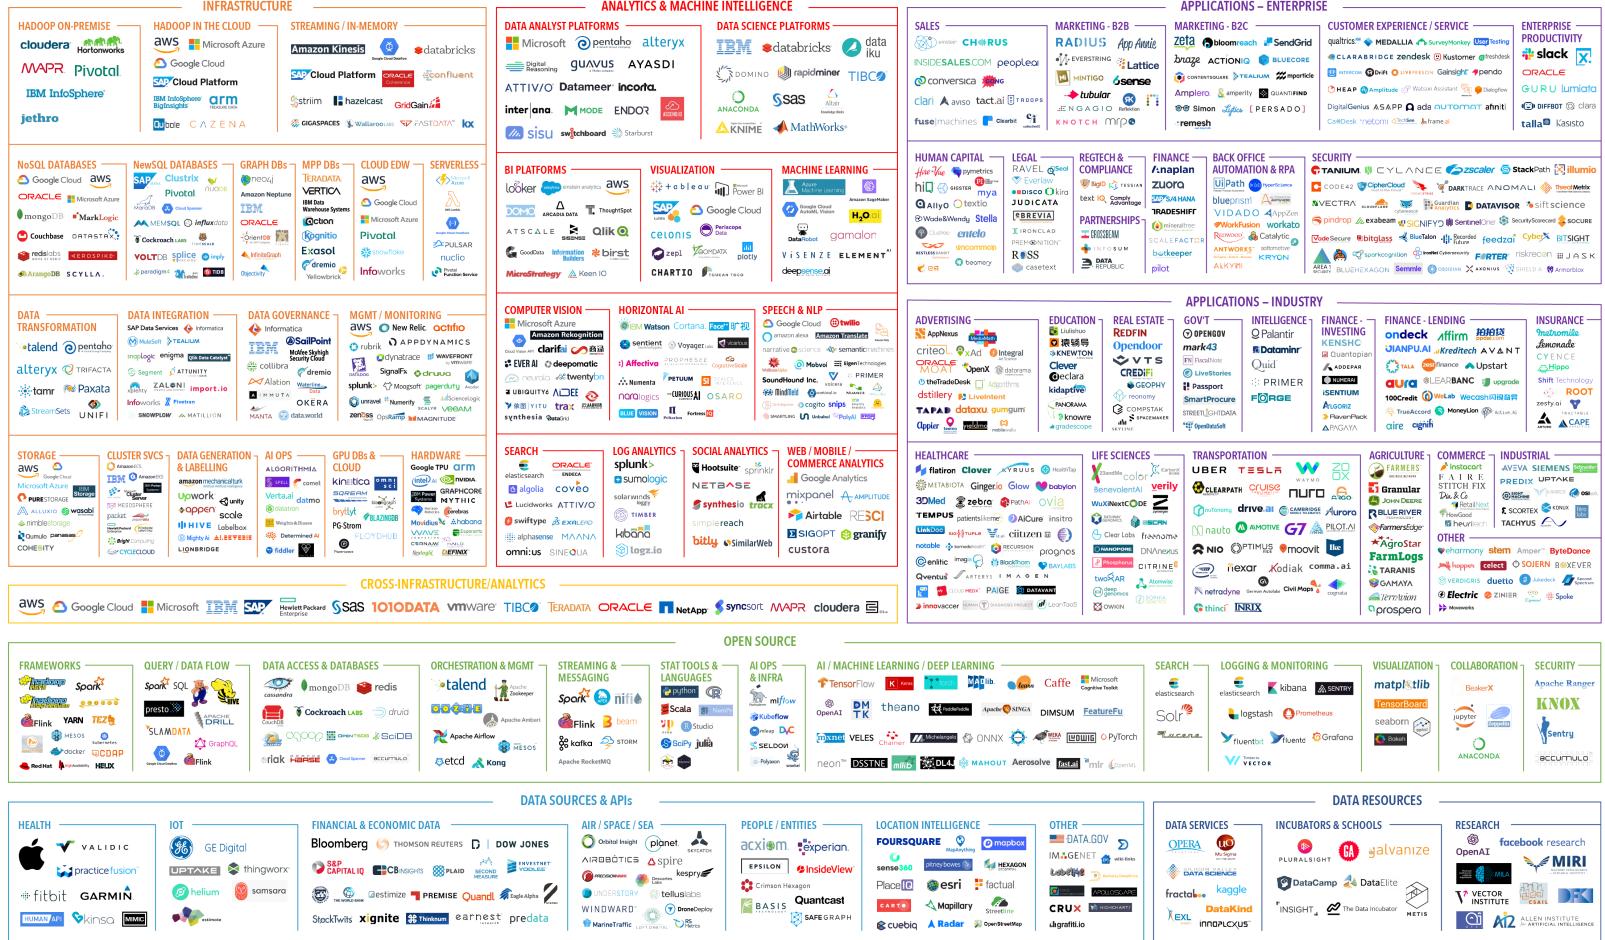
OPEN SOURCE

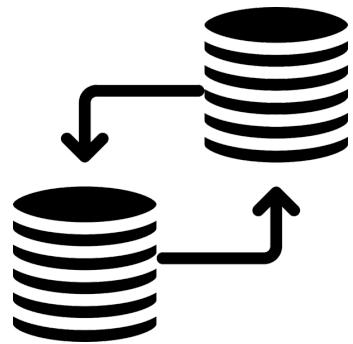


## DATA SOURCES & APIs



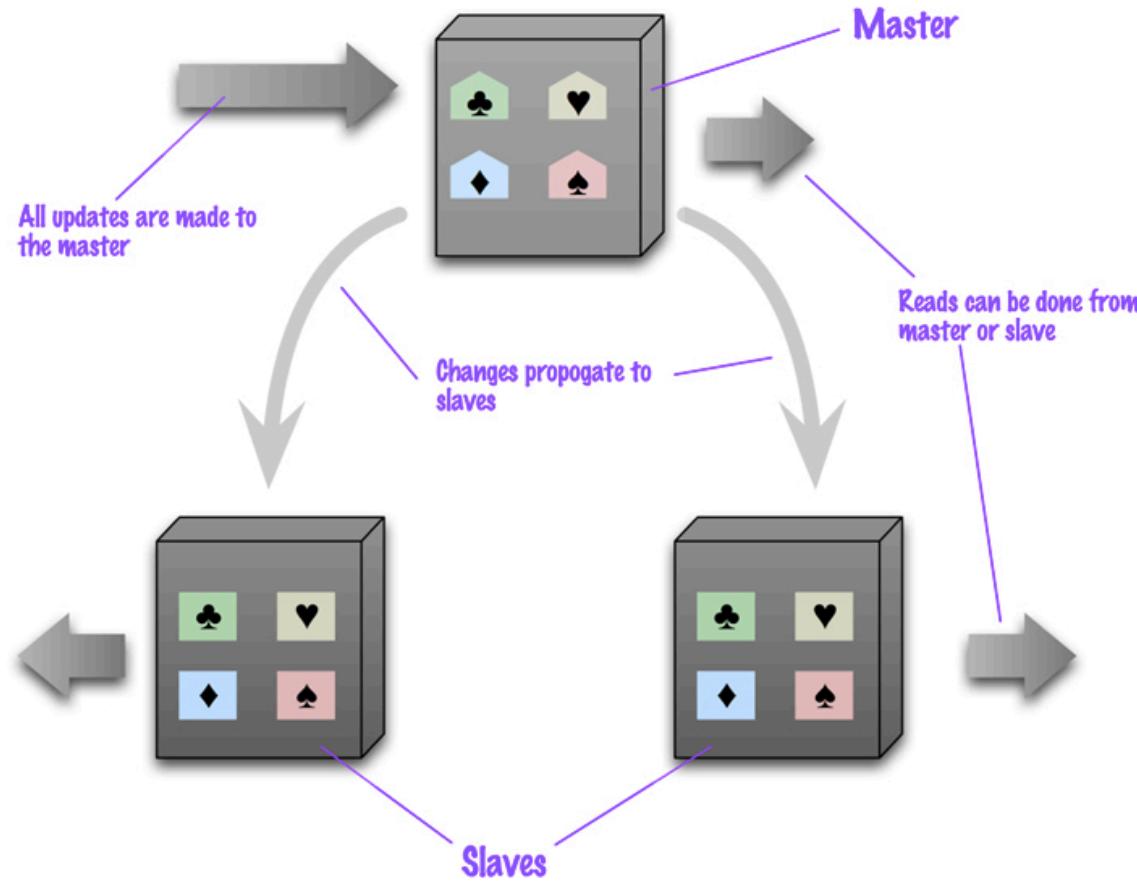
DATA & AI LANDSCAPE 2019





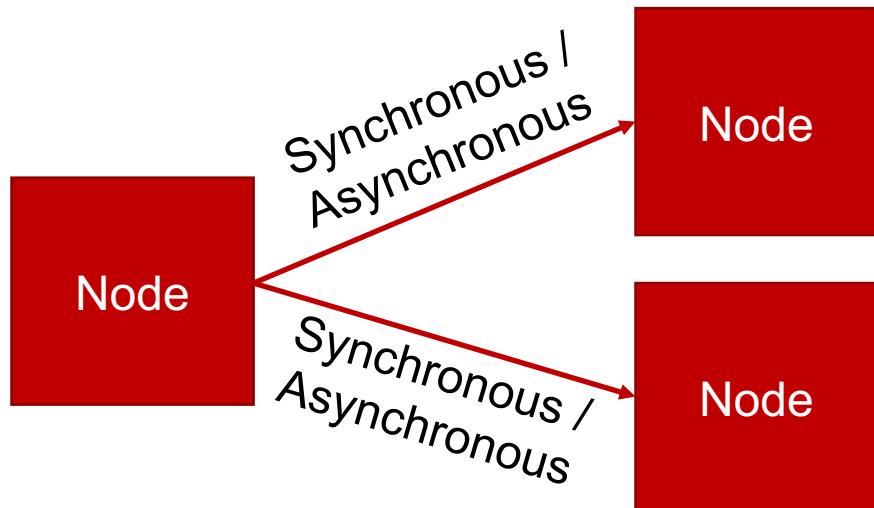
# Replication

# Replication

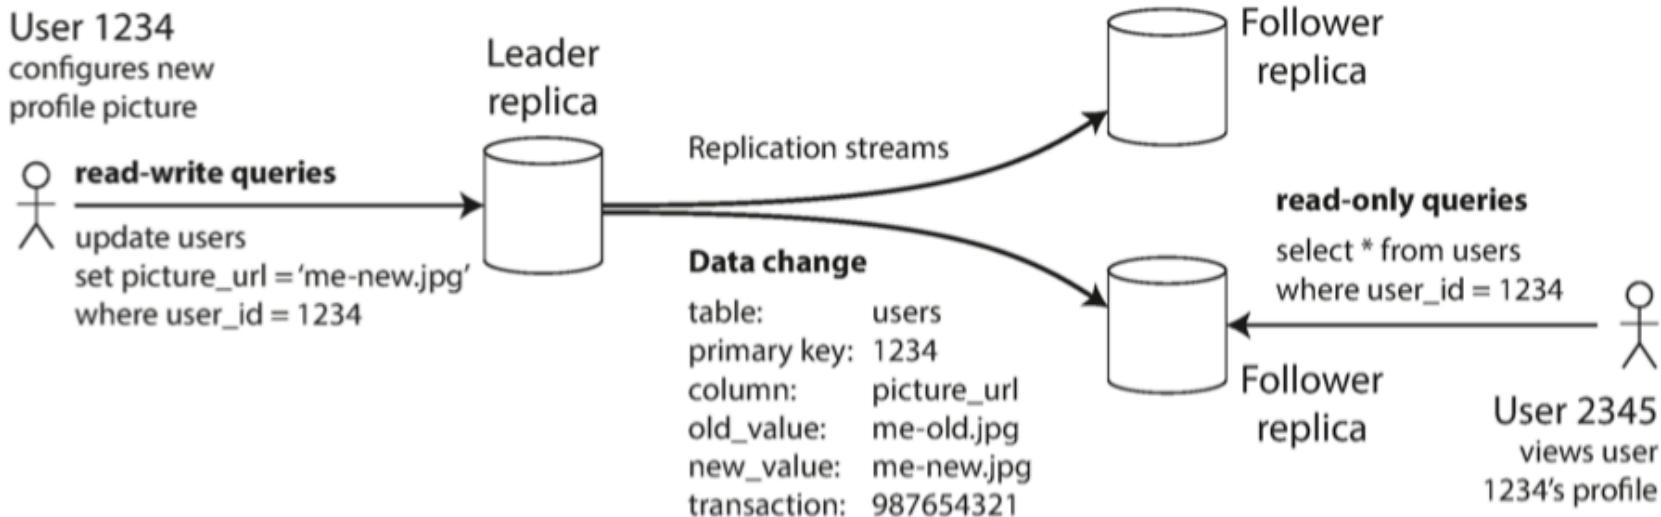


# Replication

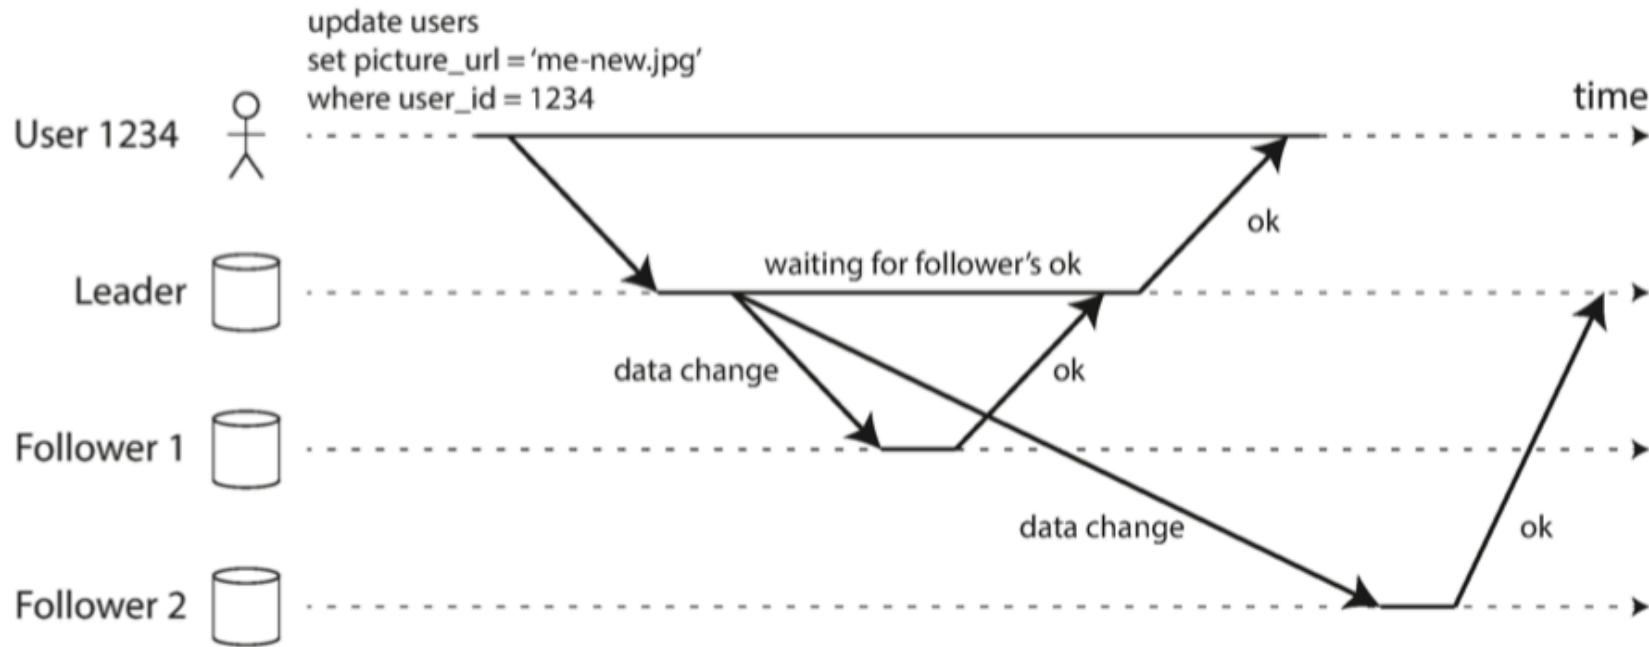
- Stores N copies of each data item
- Consistency model: synchronous vs. asynchronous
- Coordination: Single-leader, Multi-Leader, Leaderless



# Leader-based (Master–Slave) Replication



# Leader-based replication with one synchronous and one asynchronous follower



# Three main approaches to replication

## ■ Single-leader replication

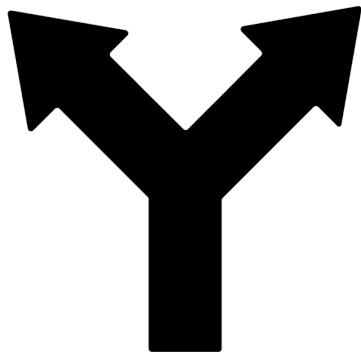
- Clients send all **writes** to a **single node** (the **leader**), which sends a stream of data change events to the other replicas (**followers**).
- **Reads** can be performed on any replica, but **reads** from followers might be **stale**.

## ■ Multi-leader replication

- Clients send each **write** to one of **several leader nodes**, any of which can accept writes.
- The leaders send streams of data change events to each other and to any follower nodes.

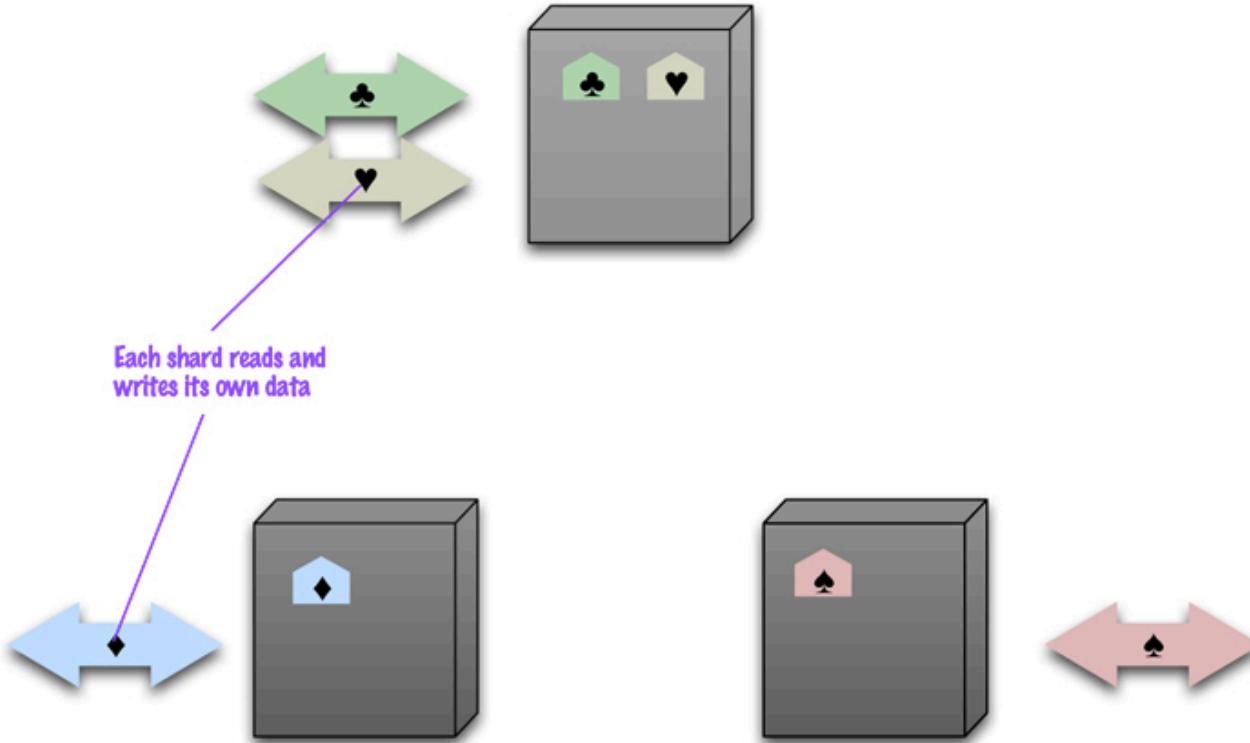
## ■ Leaderless replication

- Clients send each **write** to **several nodes**, and **read** from **several nodes** in parallel in order to detect and correct nodes with stale data.



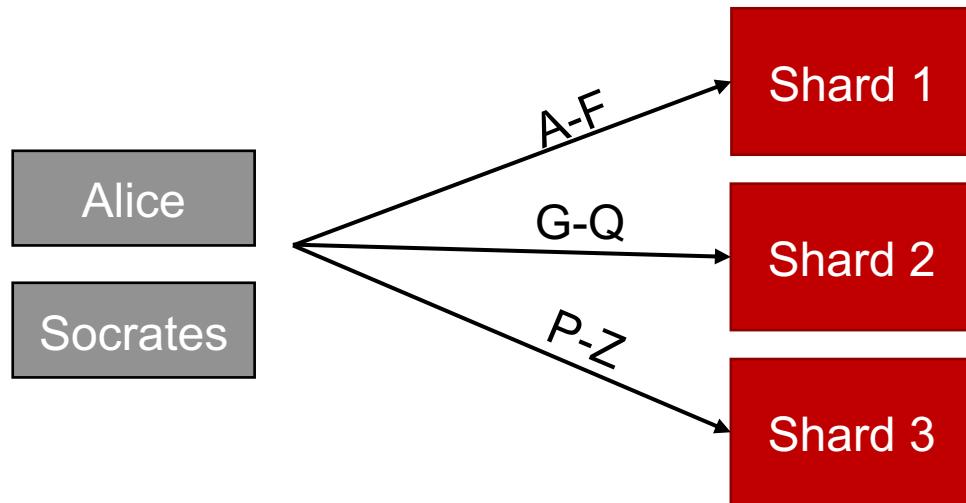
# Partitioning

# Partitioning (Sharding)

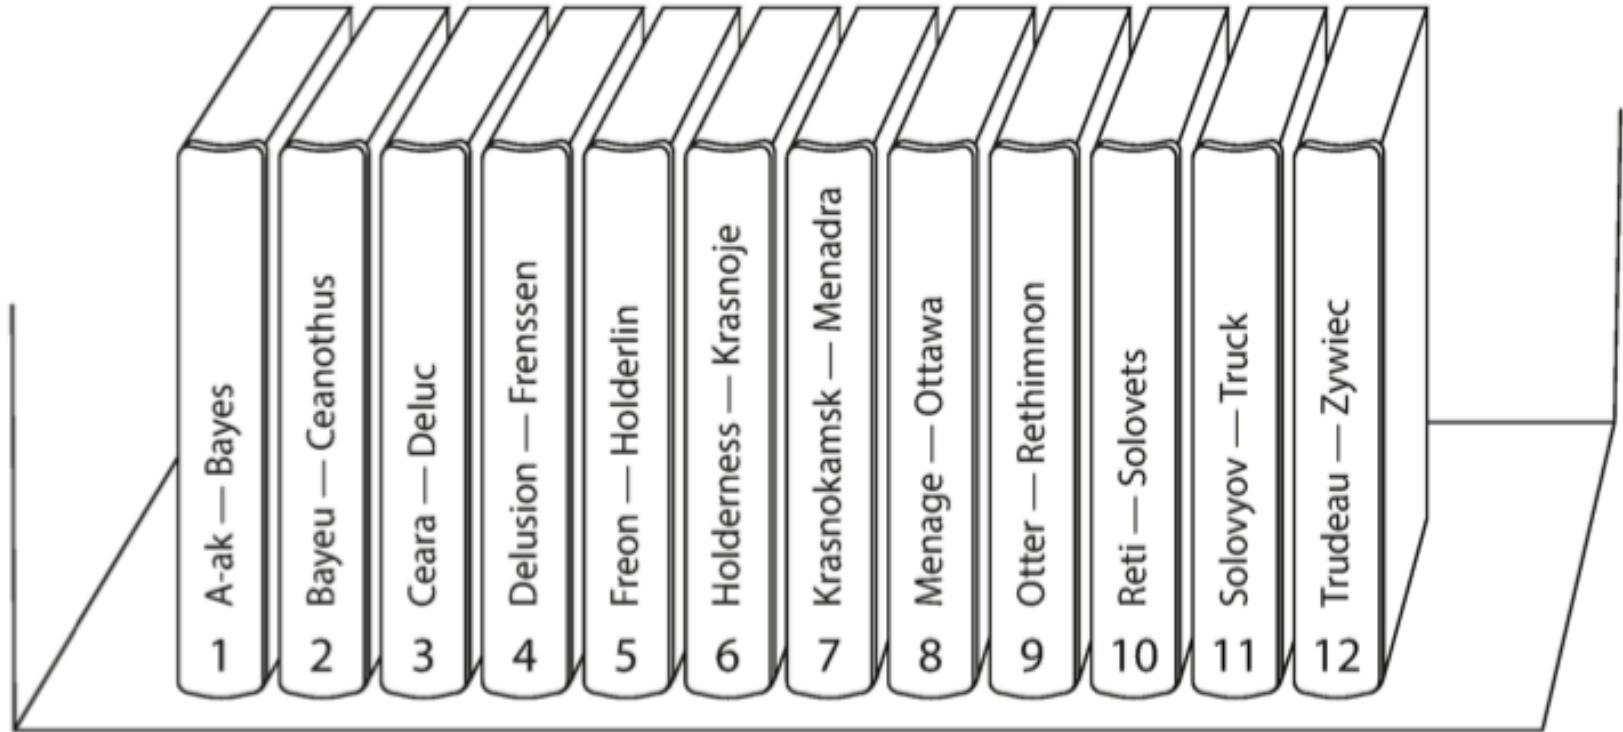


# Partitioning (Sharding)

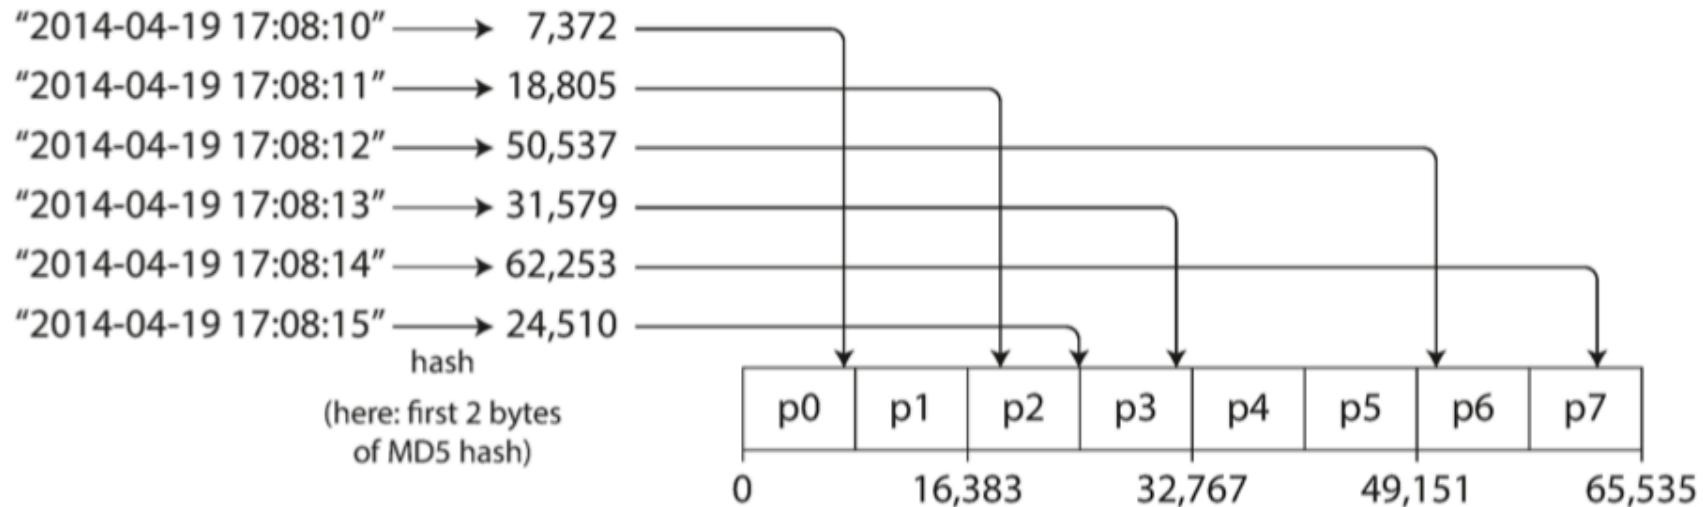
- Horizontal distribution of data over nodes
- Partitioning strategies: Hash-based vs. Range-based
- Difficulty: Multi-Shard-Operations (join, aggregation)



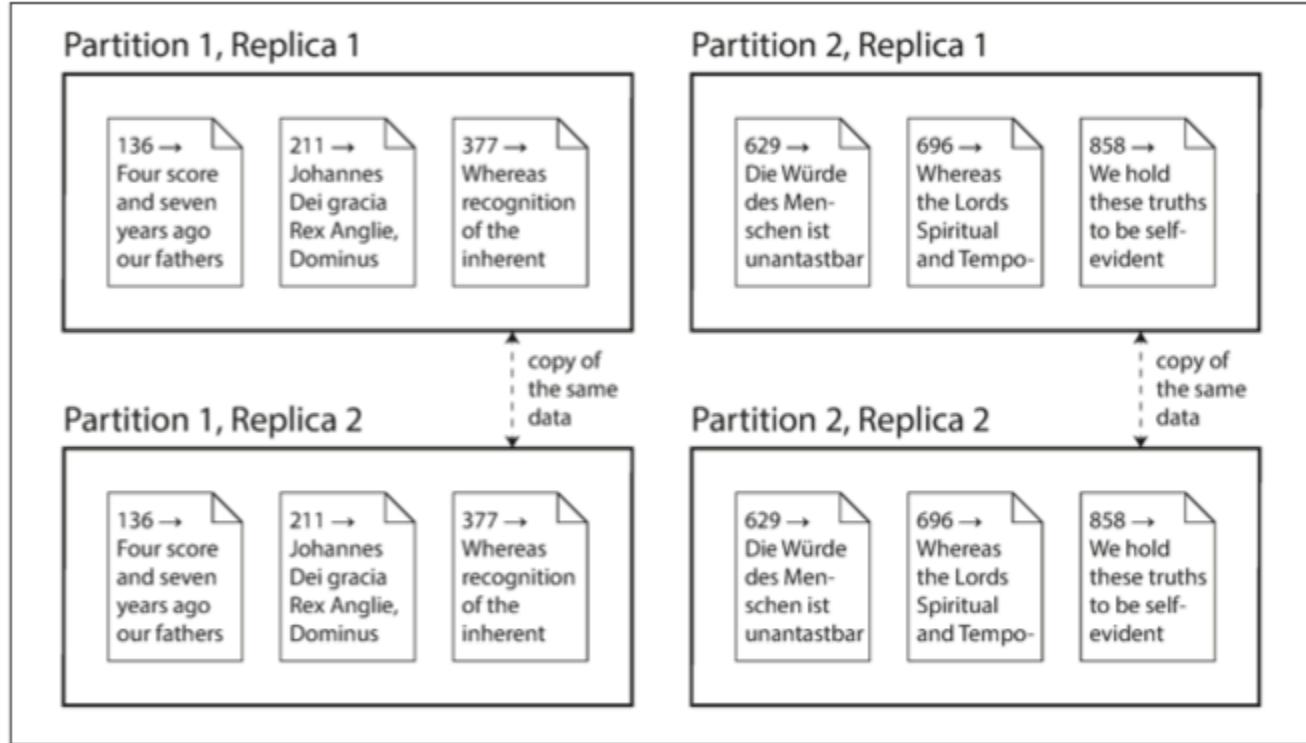
# A print encyclopedia is partitioned by key range



# Partitioning by hash of key



# Combining Replication and Partitioning



# Impact of Replication and Partitioning on

- Reliability
  - (Fault tolerance / high availability)
    - Keeping the system running, even when one machine (or several machines, or an entire datacenter) goes down
    - Disconnected operation
    - Allowing an application to continue working when there is a network interruption
- Scalability
  - Being able to handle a higher volume of reads than a single machine could handle, by performing reads on replicas or partitions
  - Storing more data than possible on a single machine by partitioning the data to more nodes
- Latency
  - Placing data geographically close to users, so that users can interact with it faster
- Maintainability
  - Difficult to operate
  - More complicated
  - Sometimes automatic failover possible
  - Problems of Distributed Consistency

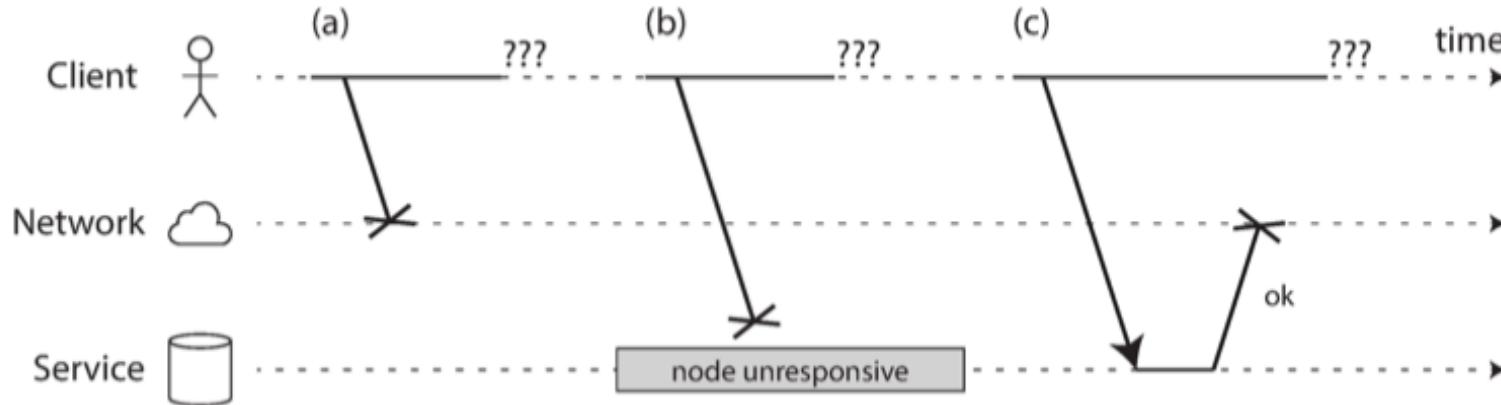


# Problems of Distributed Systems

# Problems of Distributed Systems

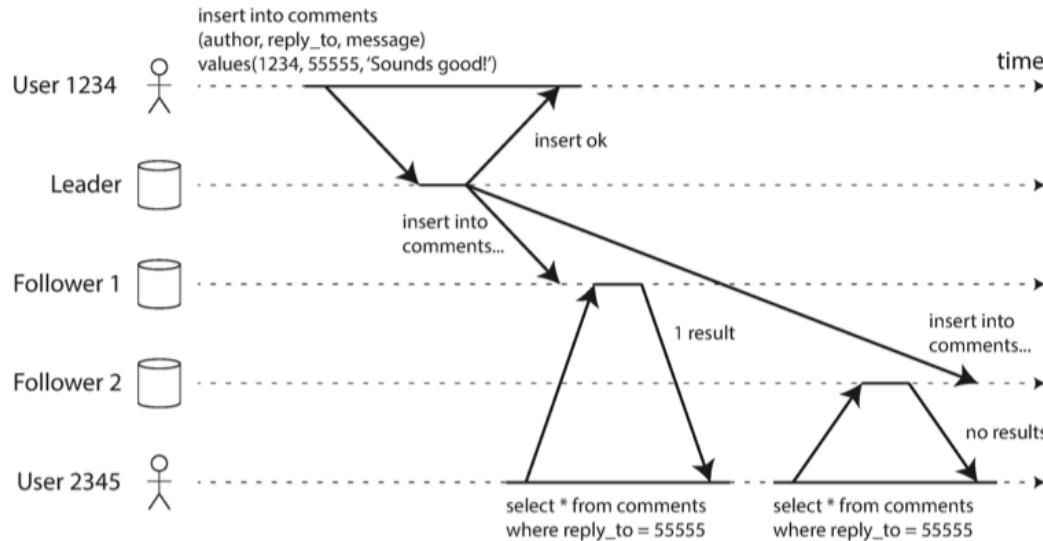
If you send a request and don't get a response, it's not possible to distinguish whether

- (a) the request was lost,
- (b) the remote node is down, or
- (c) the response was lost.



# Problems of Distributed Systems

- A user first reads from a fresh replica, then from a stale replica.
- Time appears to go backward.
- To prevent this anomaly, we need monotonic reads.



# Distributed Consistency Models

## ■ **Read-after-write consistency**

- Users should always see data that they submitted themselves.

## ■ **Monotonic reads**

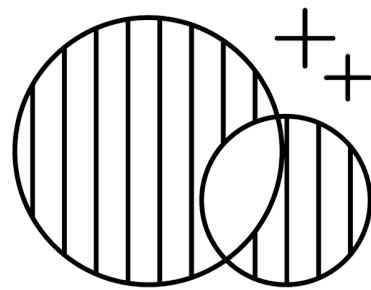
- After users have seen the data at one point in time, they shouldn't later see the data from some earlier point in time.

## ■ **Consistent prefix reads**

- Users should see the data in a state that makes causal sense: for example, seeing a question and its reply in the correct order.

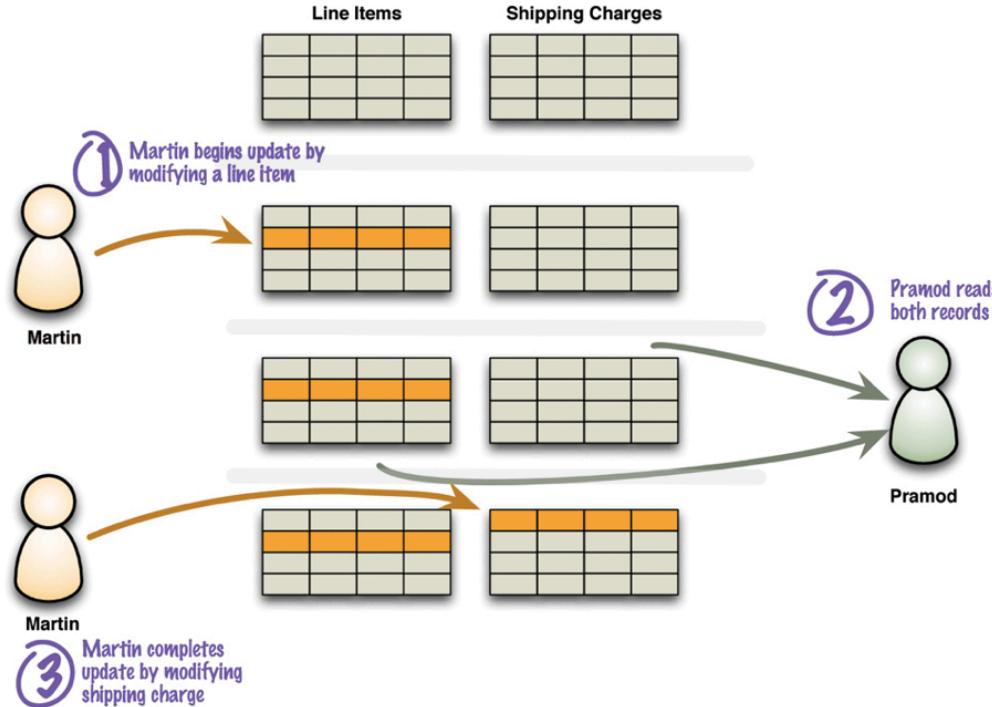
## ■ **Linearizability**

- Every user gets the same (consistent) view on the data independent with which node in a cluster he/she is connected



# CAP Theorem

# Consistency (Also known as Linearizability)



**Not the same consistency as in ACID!**

# Different Meaning of Consistency (Linearizability)

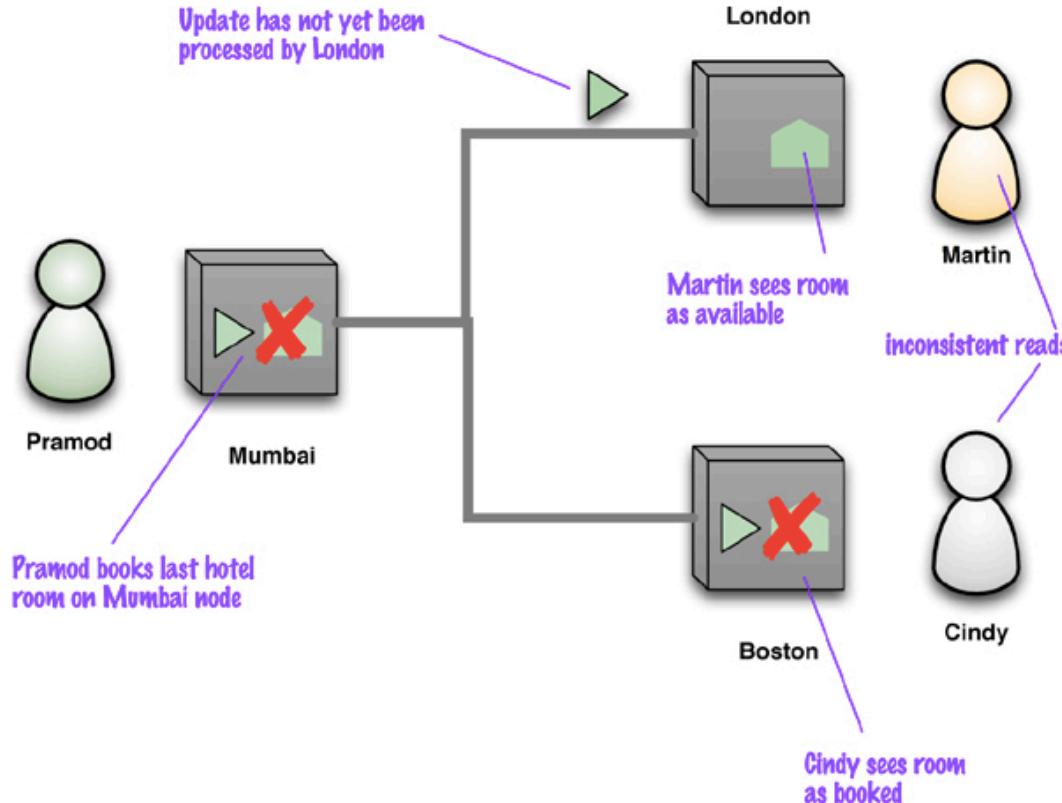
## ACID

- Any transaction will bring the database from one valid state to another
- Valid according to all defined rules and constraints

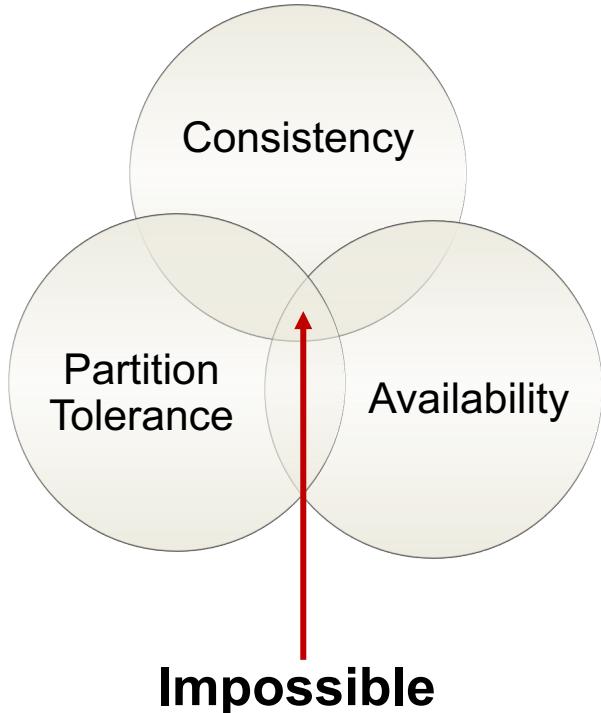
## CAP

- Every user gets the same (consistent) view on the data independent with which node in a cluster he/she is connected

# Inconsistency Problems because of Replication

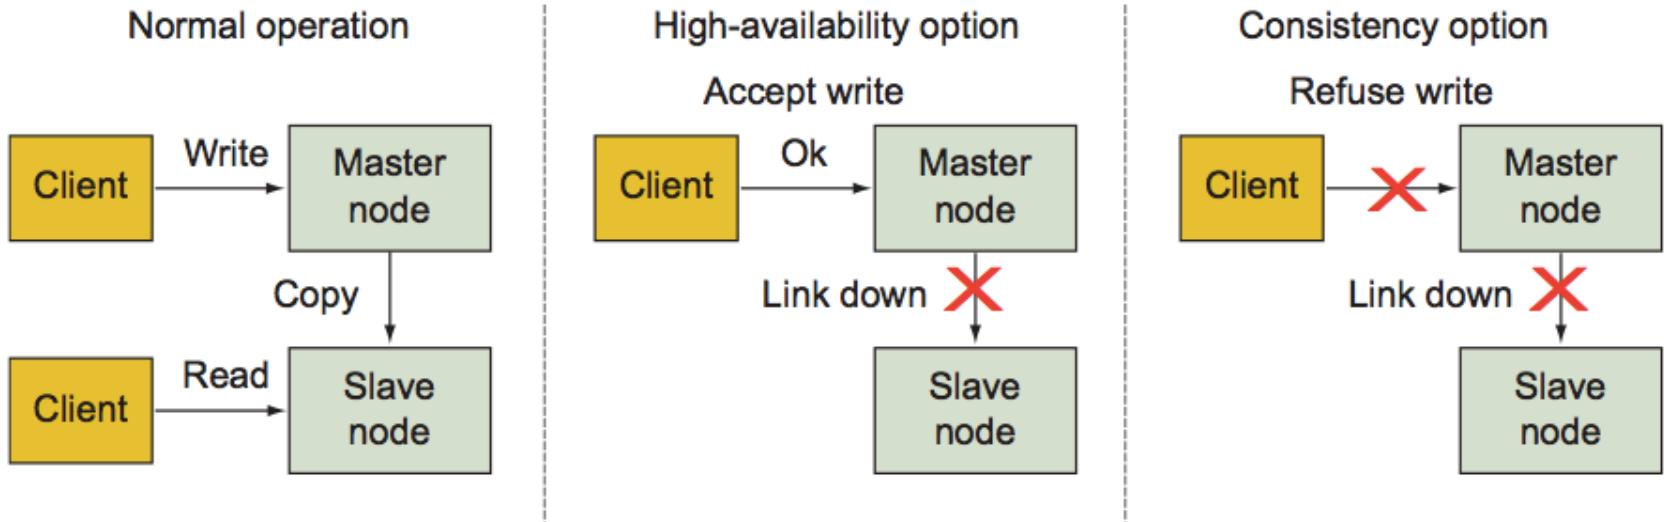


# CAP-Theorem



- Only 2 out of 3 properties are achievable at a time:
  - **Consistency:** all clients have the same view on the data
  - **Availability:** every request to a non-failed node must result in correct response
  - **Partition tolerance:** the system has to continue working, even under arbitrary network partitions

# CAP-Theorem



# CAP-Theorem

If you have...

A single processor or  
many processors on a  
working network

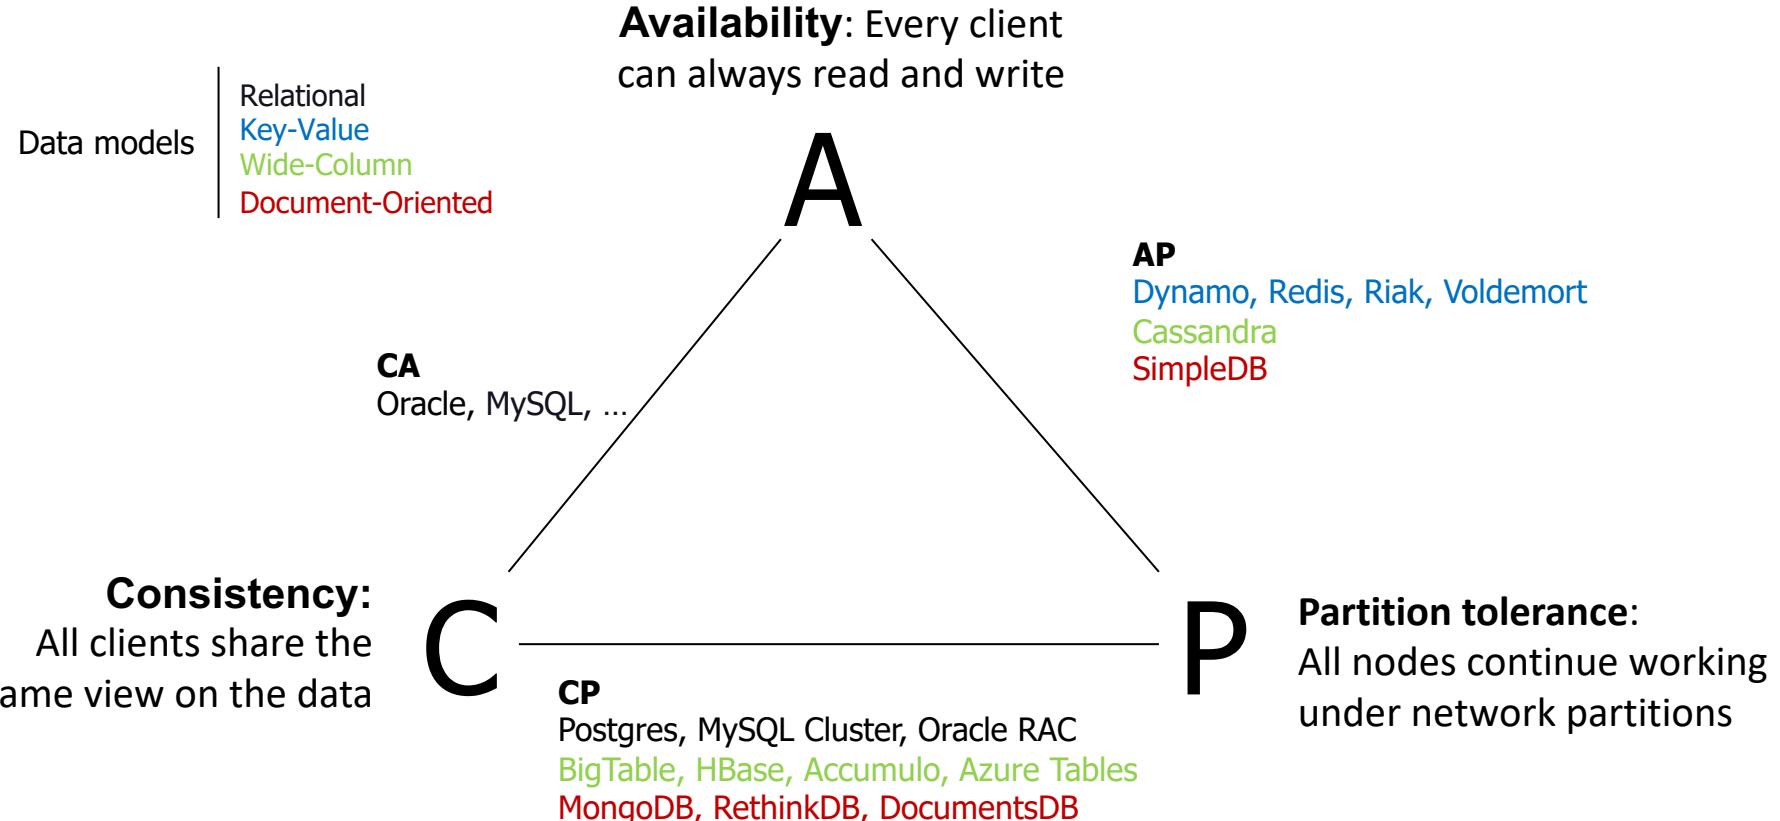
Many processors and  
network failures

then you get...

Consistency AND availability

Each transaction can select between  
consistency OR availability depending  
on the context

# CAP Triangle



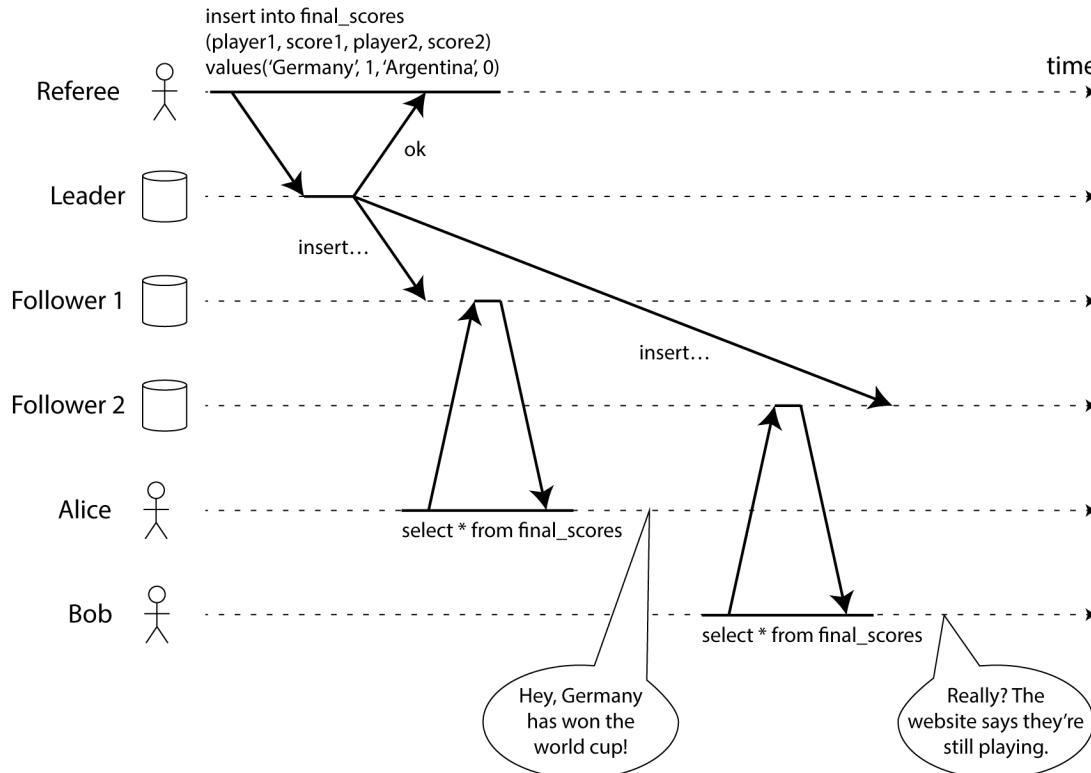
# Problematic and Unrealistic Definition of CAP Theorem

- Availability
  - Instance available. No chance of waiting for self-repair or delay
- Network Partition
  - Could be arbitrary long!
- Consistency
  - Very strong type of consistency → Linearizability
- What about Node failure?
- What about Latency?

# Consistency (Linearizability)

- „If operation B started after operation A successfully completed,
- then operation B must see the system in the same state as it was on completion of operation A, or a newer state.”

# Consistency (Linearizability)



# PACELC – Alternative for CAP

Partition: Availability vs. Consistency, Else: Latency vs. Consistency

**if there is a partition (P)**

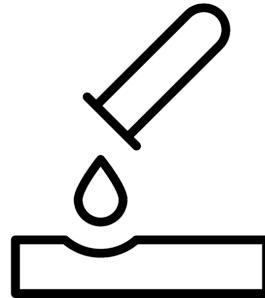
- how does the system trade off availability and consistency (A and C)

**else (E)**

- when the system is running normally in the absence of partitions, how does the system trade off latency (L) and consistency (C)

Classification of distributed storage systems:

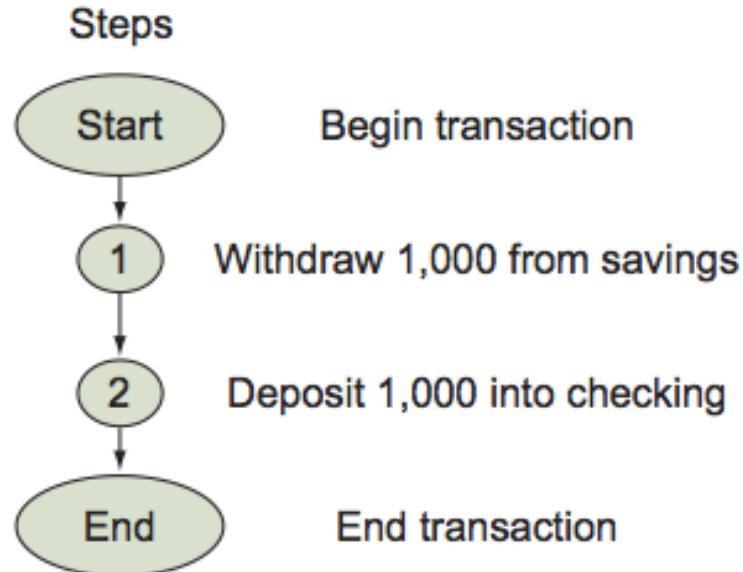
- PA/EL: Dynamo, Cassandra, Riak
- PC/EC: ACID Databases (like PostgreSQL), Big Table, Hbase
- PA/EC: MongoDB
- PC/EL: PNUTS



# ACID vs. BASE

# Transaction

From account:	Savings	▼
To account:	Checking	▼
Amount:	1,000.00	
<b>Transfer</b>		



# ACID Properties of a Transaction

- Atomicity
  - each transaction is "all or nothing": if one part of the transaction fails, the entire transaction fails, and the database state is left unchanged.
- Consistency
  - any transaction will bring the database from one valid state to another
  - valid according to all defined rules and constraints
- Isolation
  - concurrent execution of transactions results in a system state that would be obtained if transactions were executed serially, i.e., one after the other
- Durability
  - once a transaction has been committed, it will remain so, even in the event of power loss, crashes, or errors

# BASE

- Basic Availability
- Soft-state
- Eventual consistency

# ACID vs. BASE

Vs.

Acid

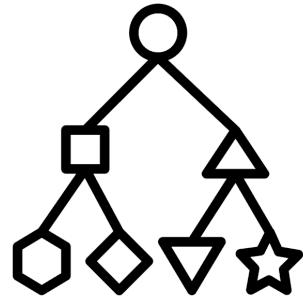
- Get transaction details right
- Block any reports while you are working
- Be pessimistic: anything might go wrong!
- Detailed testing and failure mode analysis
- Lots of locks and unlocks



Base

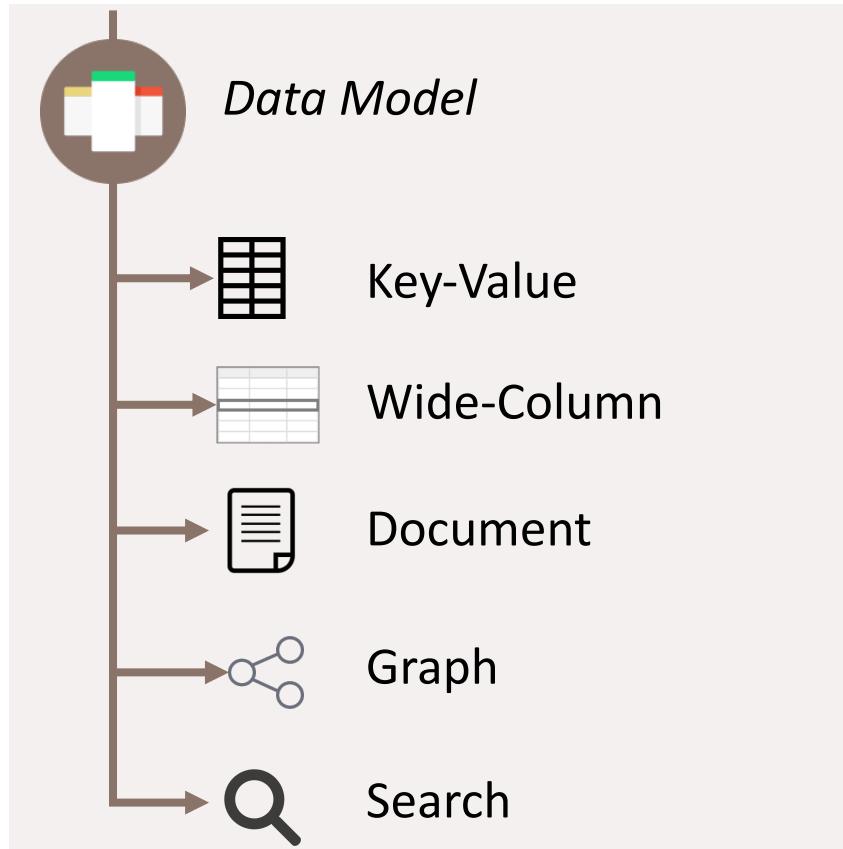
- Never block a write
- Focus on throughput, not consistency
- Be optimistic: if one service fails it will eventually get caught up
- Some reports may be inconsistent for a while, but don't worry
- Keep things simple and avoid locks





# NoSQL Classification

# NoSQL System Classification



# Key-Value Stores

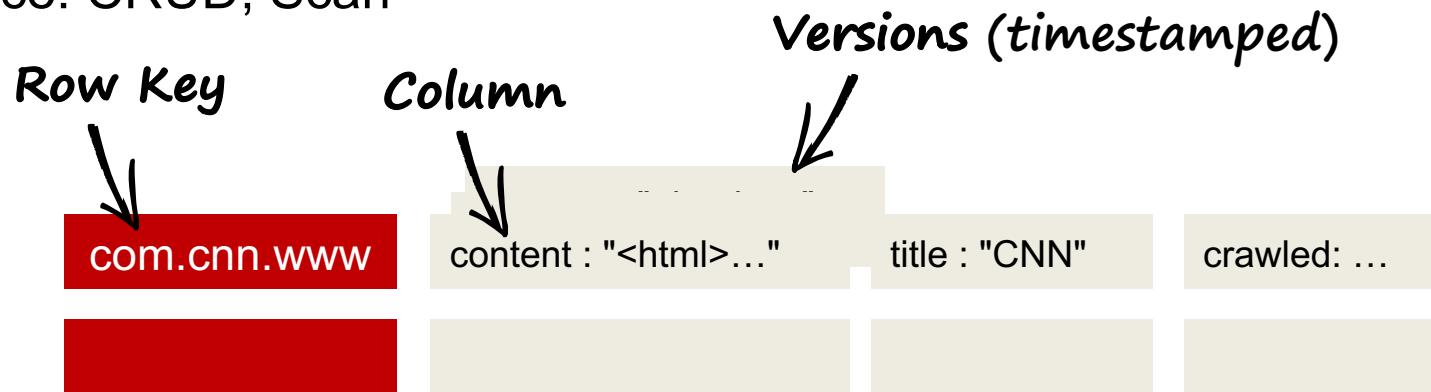
- Data model: (key) -> value
- Interface: CRUD (Create, Read, Update, Delete)



- Examples: Redis (CP), Amazon Dynamo (AP), Riak (AP)

# Wide-Column Stores

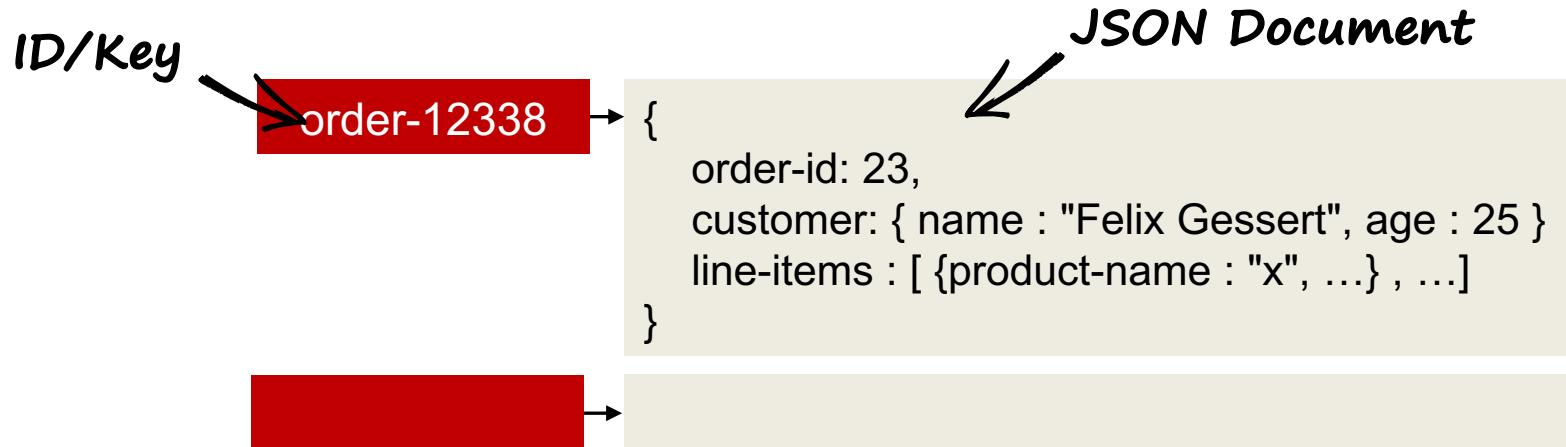
- Data model: (rowkey, column, timestamp) -> value
- Interface: CRUD, Scan



- Examples: Cassandra (AP), Google BigTable (CP), HBase (CP)
- Two-dimensional key-value store
- NOT equal to column-oriented storage like in SAP HANA!

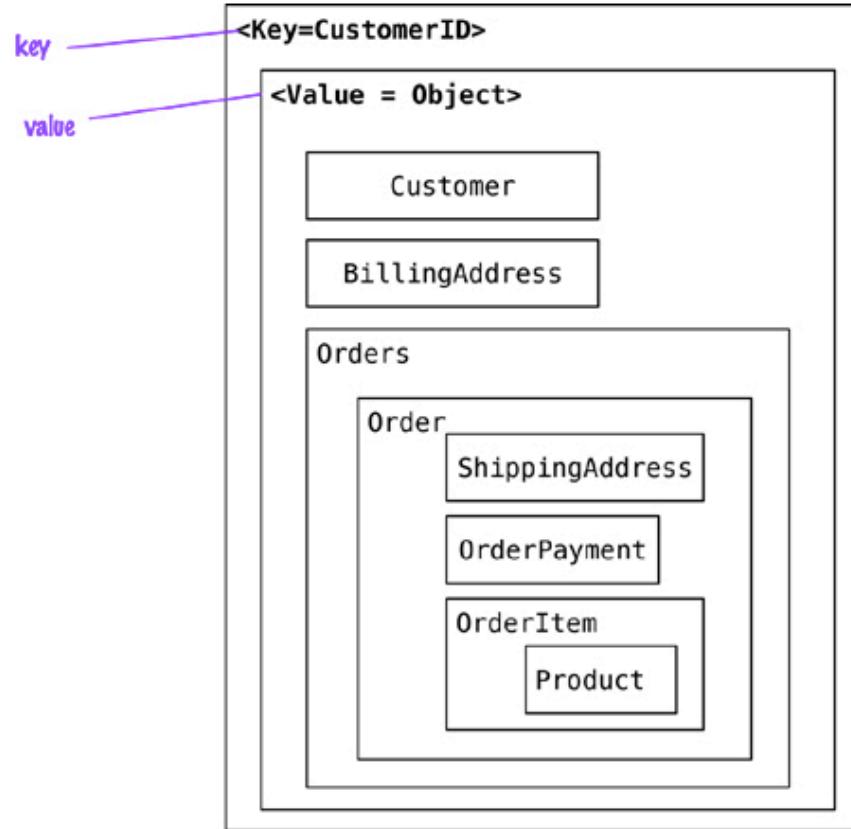
# Document Stores

- Data model: (collection, key) -> document
- Interface: CRUD, Querys, Map-Reduce



- Examples: MongoDB (CP), CouchDB (AP), RethinkDB (CP)

# Document Stores

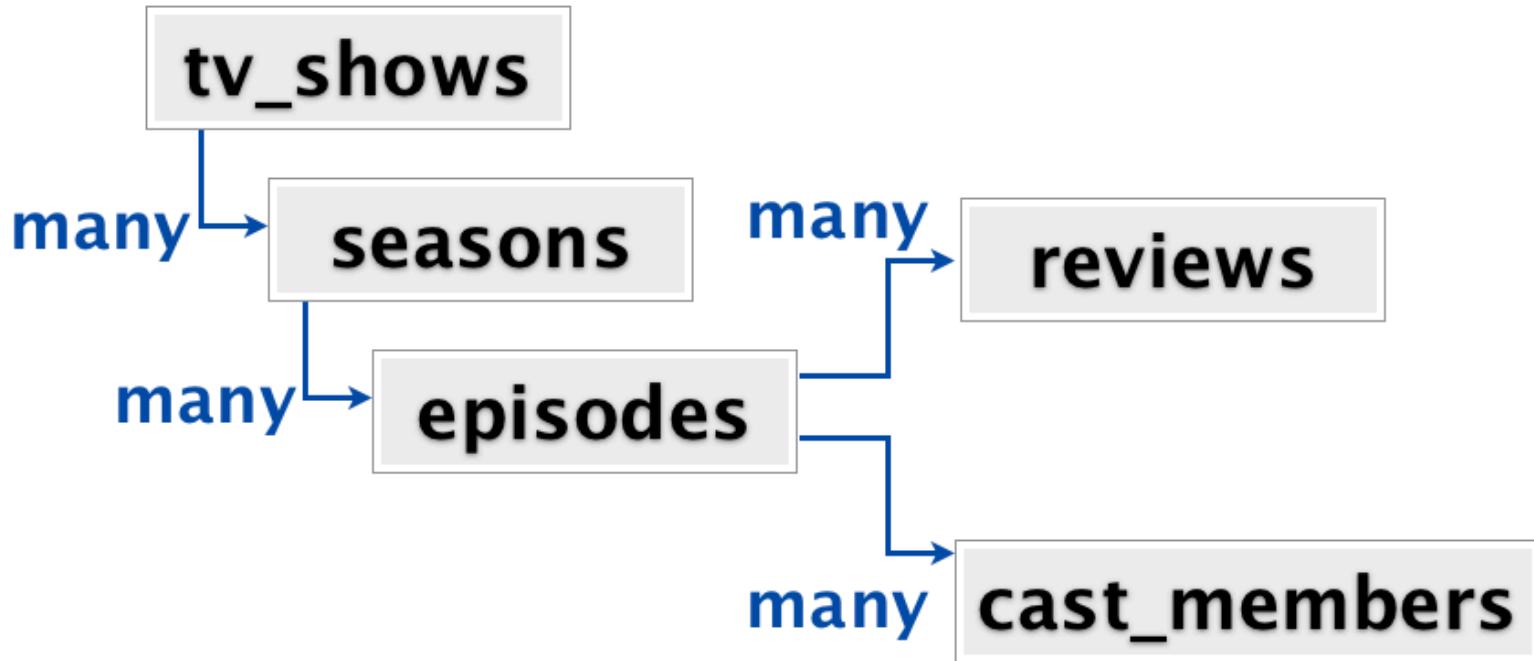


# JSON

```
# Customer object
{
  "customerId": 1,
  "customer": {
    "name": "Martin",
    "billingAddress": [{"city": "Chicago"}],
    "payment": [{"type": "debit", "ccinfo": "1000-1000-1000-1000"}],
    "orders": [{"orderId": 99}]
  }
}

# Order object
{
  "customerId": 1,
  "orderId": 99,
  "order": {
    "orderDate": "Nov-20-2011",
    "orderItems": [{"productId": 27, "price": 32.45}],
    "orderPayment": [{"ccinfo": "1000-1000-1000-1000",
                     "txnId": "abelif879rft"}],
    "shippingAddress": {"city": "Chicago"}
  }
}
```

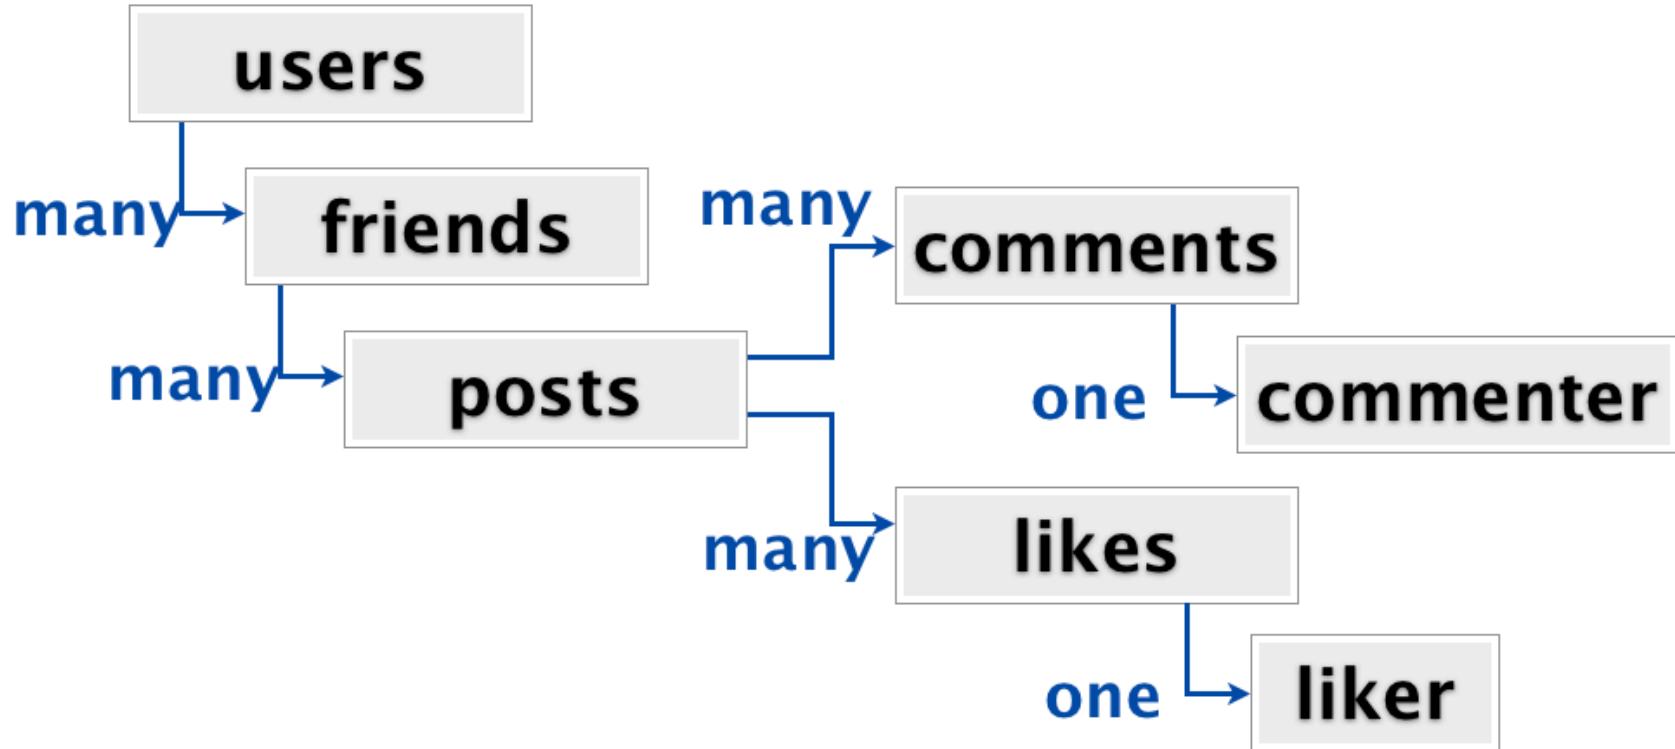
# Not all Tree Structures should be saved in Document Databases



# Not all Tree Structures should be saved in Document Databases

```
{title: 'Babylon 5',
  seasons: [
    {season_number: '1',
      episodes: [
        {ordinal_within_season: '1',
          title: 'Midnight on the Firing Line'
          reviews: [...],
          cast_members: [...]
        }
      ]
    }
  ]
}
```

# Not all Tree Structures should be saved in Document Databases



# MongoDB (CP)

- From humongous  $\approx$  gigantic
- Schema-free document database with tunable consistency
- Allows complex queries and indexing
- Sharding (either range- or hash-based)
- Replication (either synchronous or asynchronous)
- Storage Management:
  - Write-ahead logging for redos (journaling)
  - Storage Engines: memory-mapped files, in-memory, Log-structured merge trees (WiredTiger), ...

# Basics of MongoDB

```
> mongod &
> mongo imdb
MongoDB shell version: 2.4.3
connecting to: imdb
> show collections
movies
tweets
> db.movies.findOne({title : "Iron Man 3"})
{
  title : "Iron Man 3",
  year : 2013 ,
  genre : [
    "Action",
    "Adventure",
    "Sci -Fi"],
  actors : [
    "Downey Jr., Robert",
    "Paltrow , Gwyneth",]
}
```

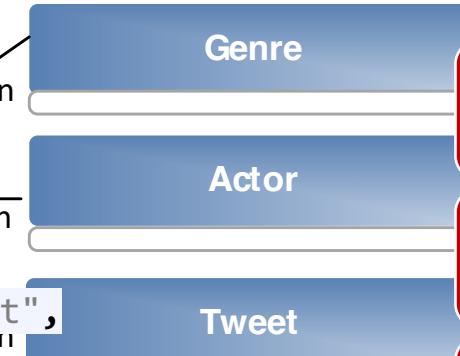
Properties

Arrays, Nesting allowed

# Data Modelling of MongoDB

```
{  
    "_id" : ObjectId("51a5d316d70beffe74ecc940")  
    title : "Iron Man 3",  
    year : 2013,  
    rating : 7.6,  
    director: "Shane Block",  
    genre : [ "Action",  
              "Adventure",  
              "Sci -Fi"],  
    actors : ["Downey Jr., Robert",  
              "Paltrow , Gwyneth"],  
    tweets : [ {  
                "user" : "Franz Kafka",  
                "text" : "#nowwatching Iron Man 3",  
                "retweet" : false,  
                "date" : ISODate("2013-05-29T13:15:51Z")  
              }]  
}
```

**Movie Document**



**Denormalisation**  
instead of joins

**Nesting** replaces 1:n  
and 1:1 relations

**Schemafreeness:**  
Attributes per document

**Unit of atomicity:**  
document

**Principles**

# Sharding and Replication of MongoDB

## Sharding:

- Sharding attribute
- Hash vs. range sharding



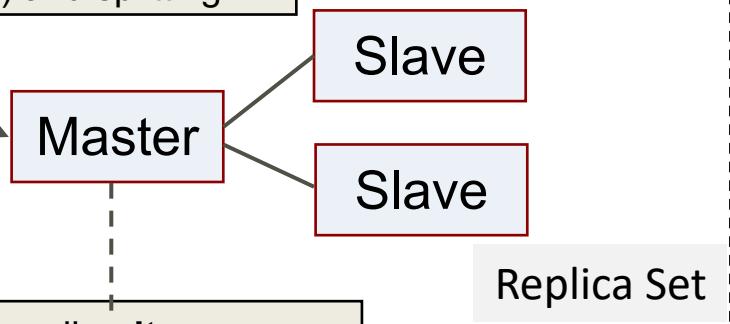
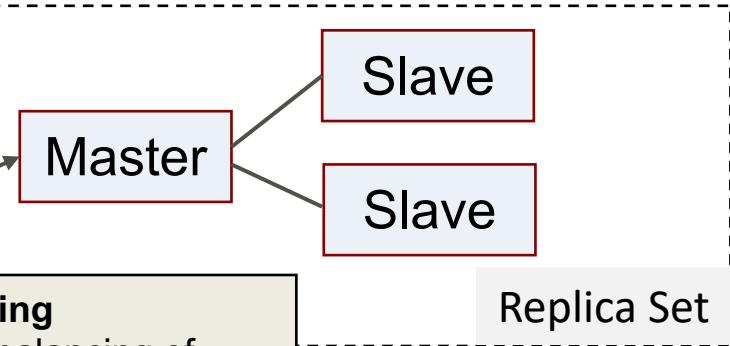
Client



Client



**Load-Balancing**  
- can trigger rebalancing of chunks (64MB) and splitting



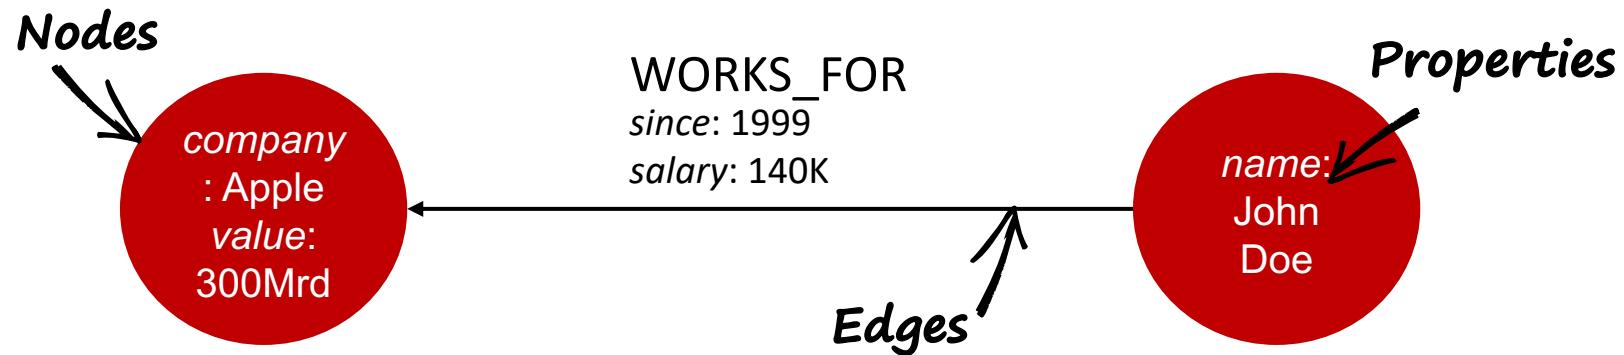
## Controls Write Concern:

- Unacknowledged, Acknowledged,*
- Journaled, Replica Acknowledged*

- Receives all **writes**  
- **Replicates** asynchronously

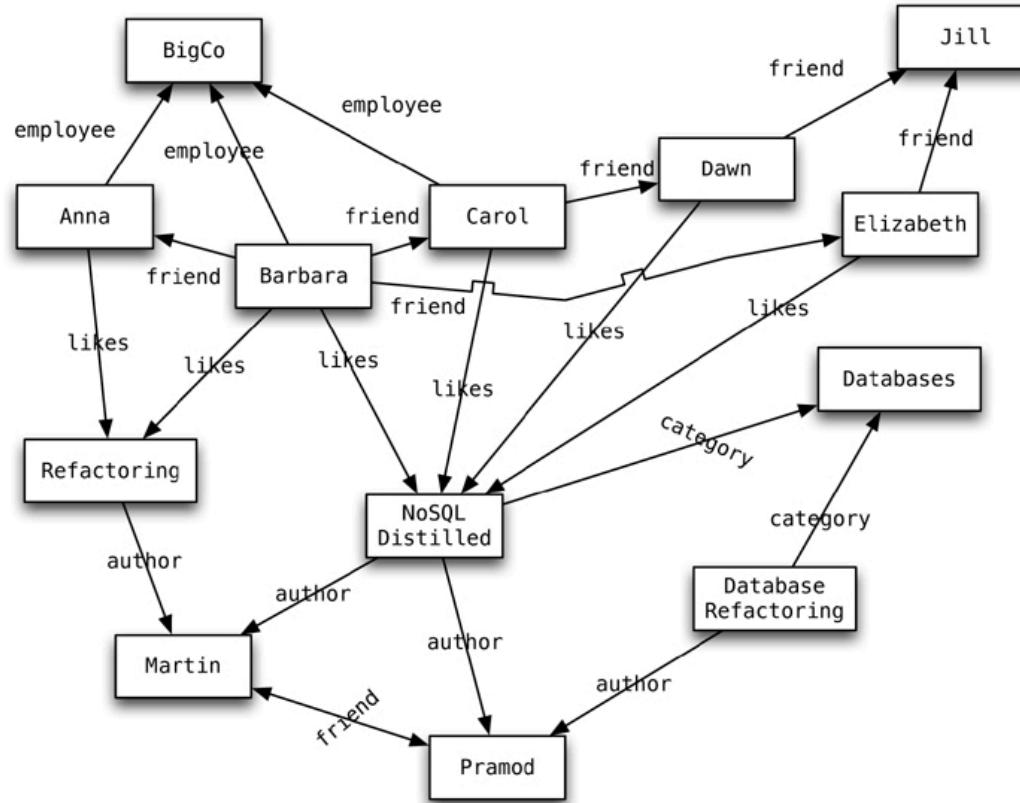
# Graph Databases

- Data model:  $G = (V, E)$ : Graph-Property Modell
- Interface: Traversal algorithms, querys, transactions

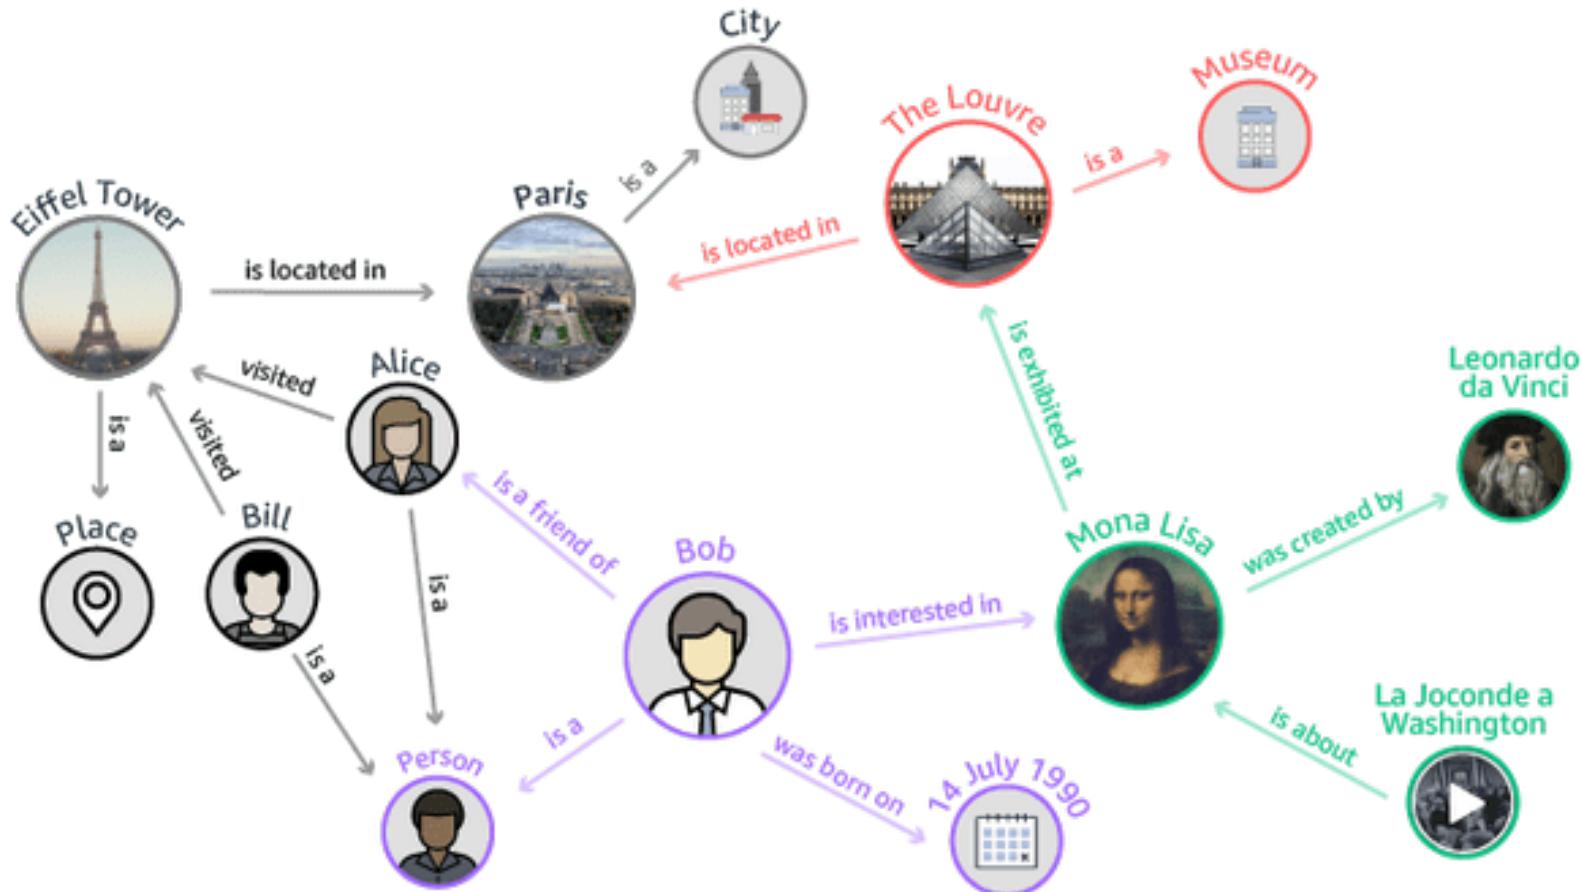


- Examples: Neo4j (CA), InfiniteGraph (CA), OrientDB (CA)

# Graph Structures

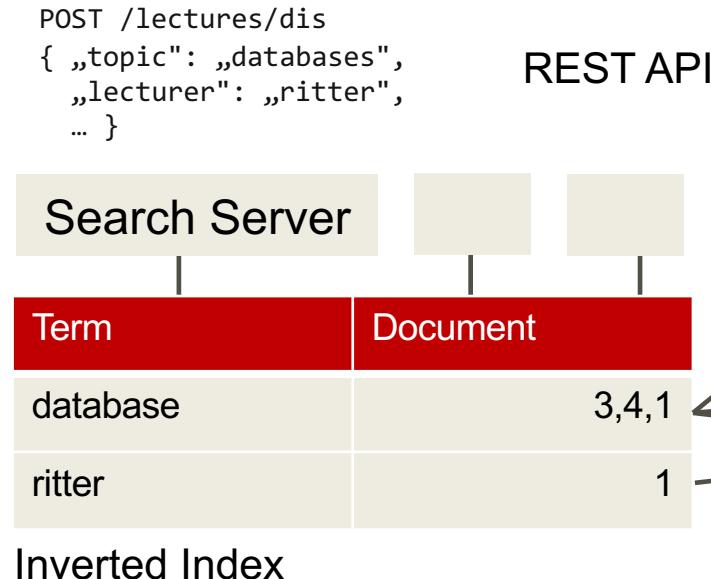


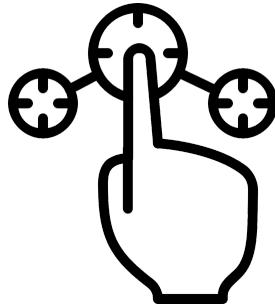
# Knowledge Graph



# Search Platforms

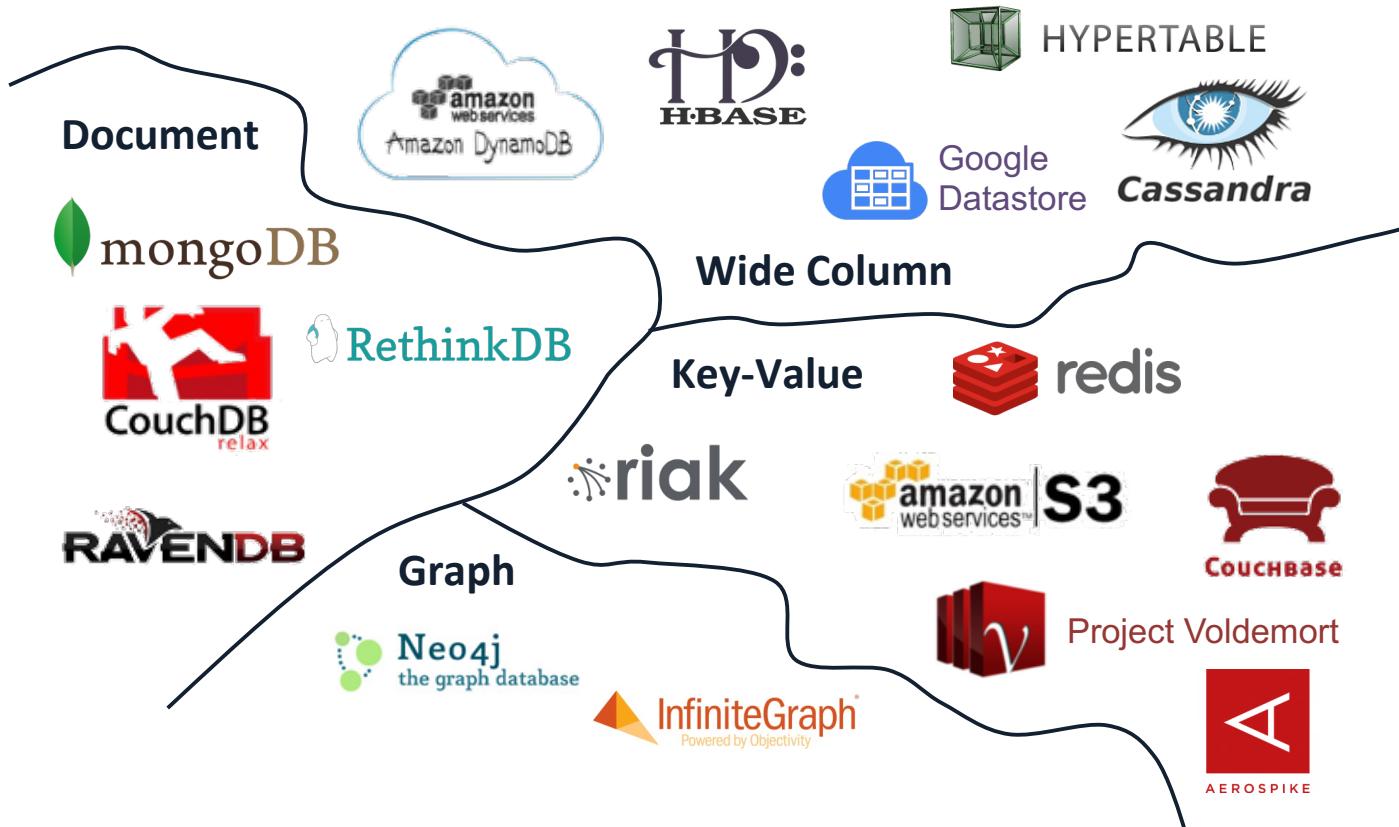
- Data model: vectorspace model, docs + metadata
- Examples: ElasticSearch, Solr





# NoSQL Selection

# NoSQL Landscape



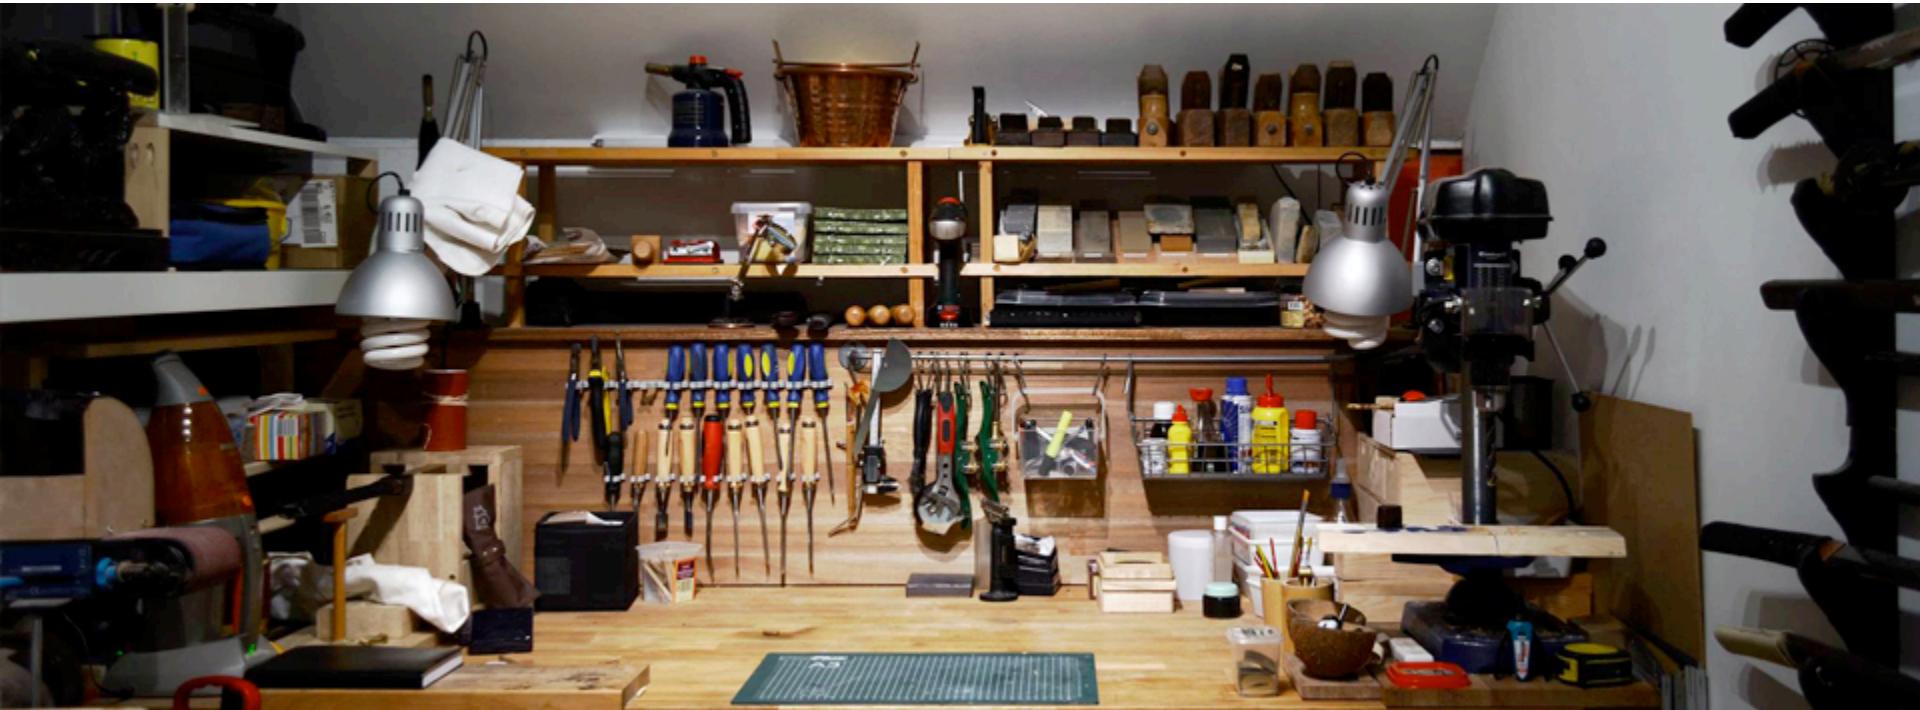
# Database Popularity May 2020

#	System	Model	Score
1.	Oracle	Relational DBMS	1345
2.	MySQL	Relational DBMS	1282
3.	MS SQL Server	Relational DBMS	1078
4.	PostgreSQL	Relational DBMS	514
5.	<b>MongoDB</b>	<b>Document store</b>	439
6.	DB2	Relational DBMS	162
7.	<b>Elasticsearch</b>	<b>Search engine</b>	149
8.	Redis	<b>Key-value store</b>	143
9.	SQLite	Relational DBMS	123
10.	MS Access	Relational DBMS	120

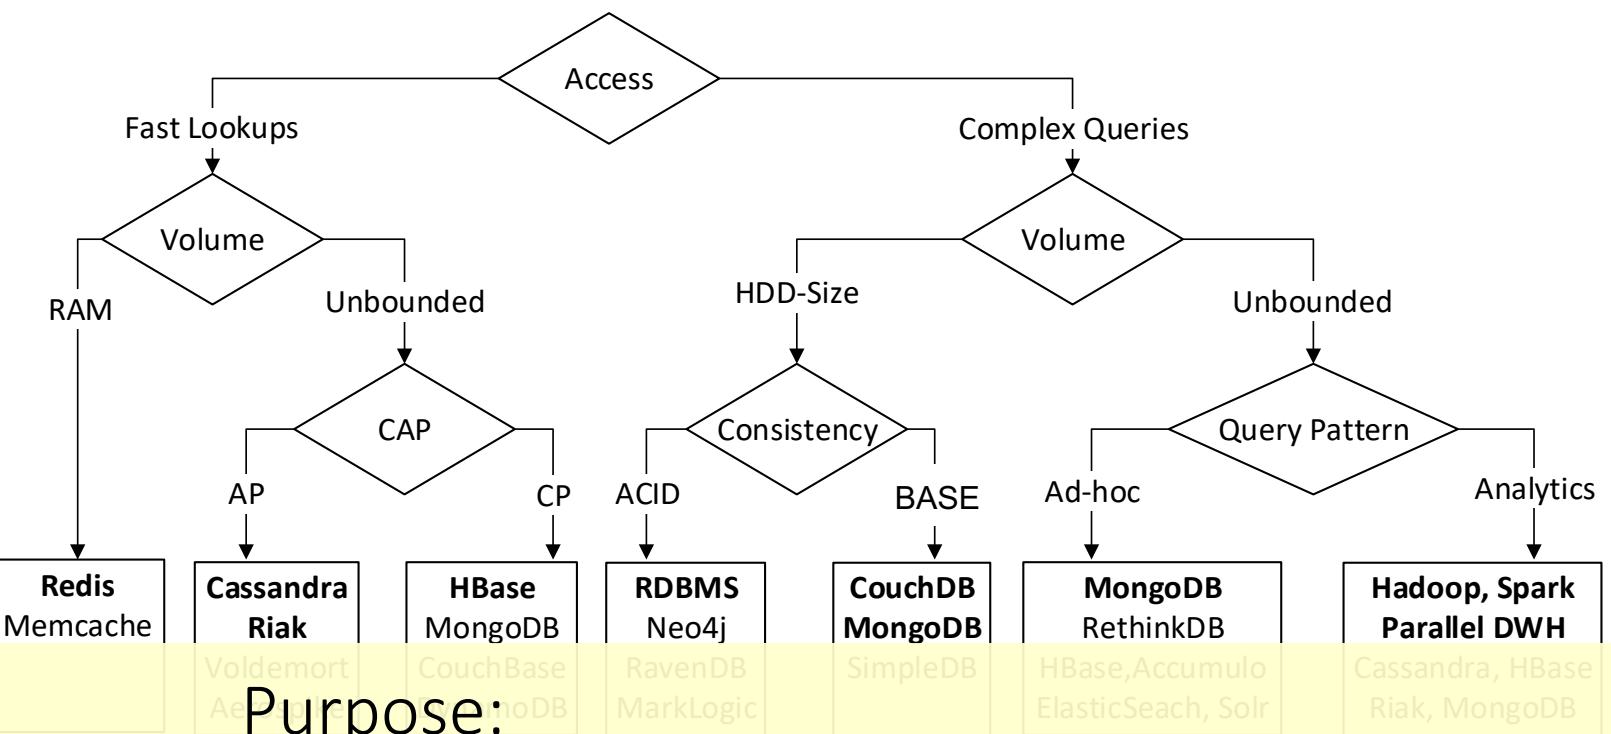
11.	<b>Cassandra</b>	<b>Wide column store</b>	119
12.	MariaDB	Relational DBMS	90
13.	<b>Splunk</b>	<b>Search engine</b>	88
14.	Hive	Relational DBMS	82
15.	Terradata	Relational DBMS	74
16.	<b>Amazon DynamoDB</b>	<b>Multi-modal store</b>	65
17.	SAP Adaptive Server	Relational DBMS	54
18.	<b>Solr</b>	<b>Search engine</b>	53
19.	FileMaker	Relational DBMS	51
20.	SAP HANA	Relational DBMS	51
21.	<b>Neo4j</b>	<b>Graph DBMS</b>	50
22.	<b>HBase</b>	<b>Wide column store</b>	50
23.	Azure SQL Database	Relational DBMS	43
24.	<b>Azure Cosmos DB</b>	<b>Multi-modal store</b>	31
25.	<b>Couchbase</b>	<b>Document store</b>	29

**Scoring:** Google/Bing results, Google Trends, Stackoverflow, job offers, LinkedIn

# Know your Tools



“Having a big toolbox is impressive. But knowing **when to use what tool** and why separates the true expert from the show-off.” *Gregor Hohpe*

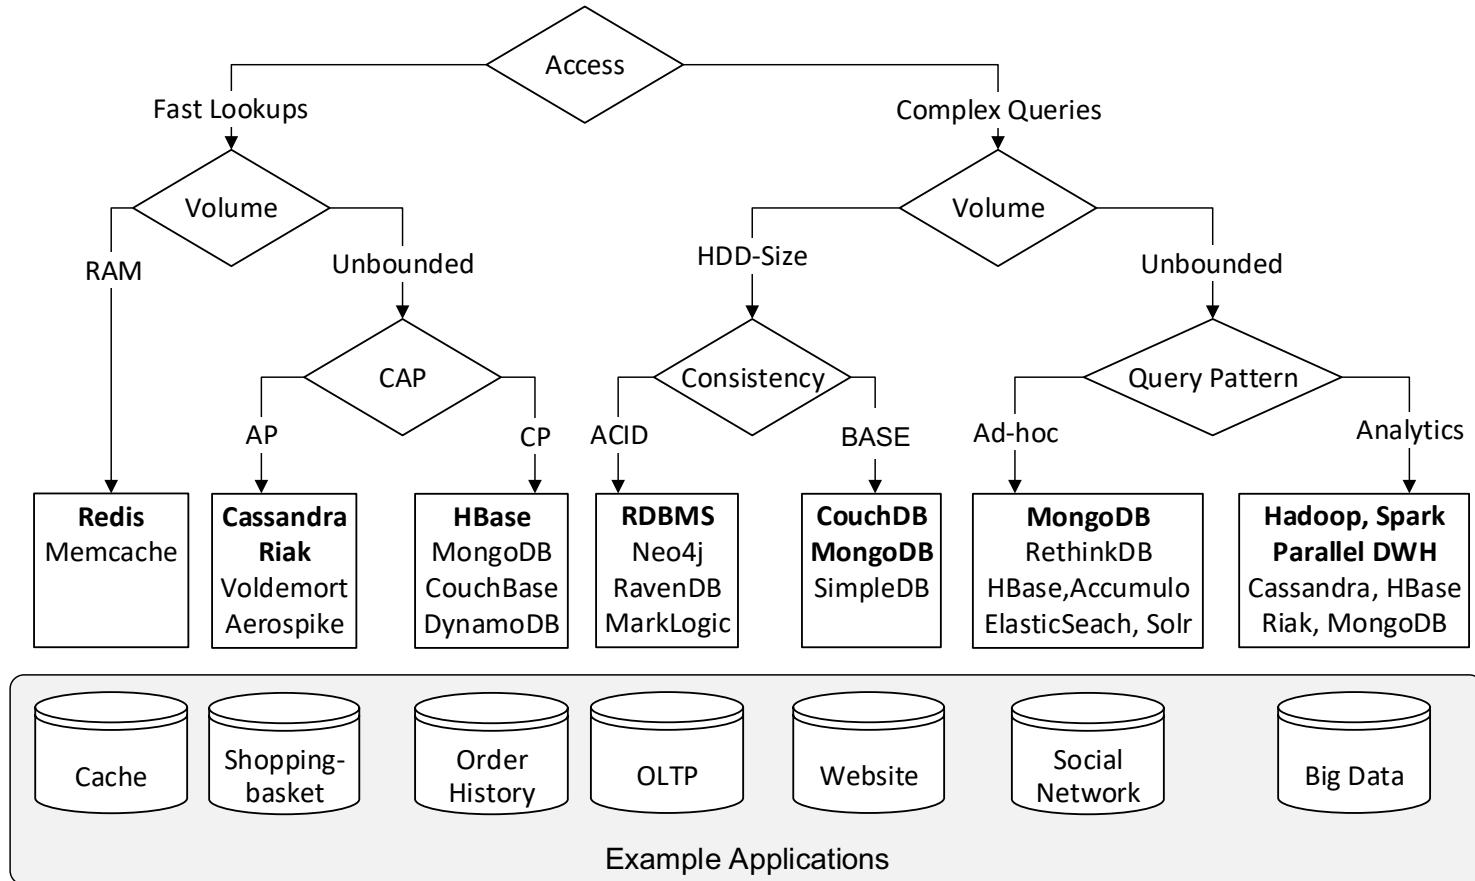


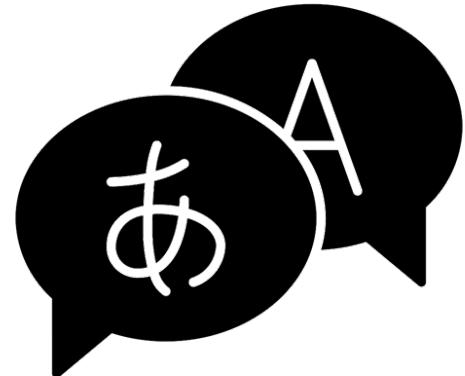
Purpose:

**Application Architects:** narrowing down the potential system candidates based on requirements

**Database Vendors/Researchers:** clear communication and design of system trade-offs

# NoSQL Decision Tree

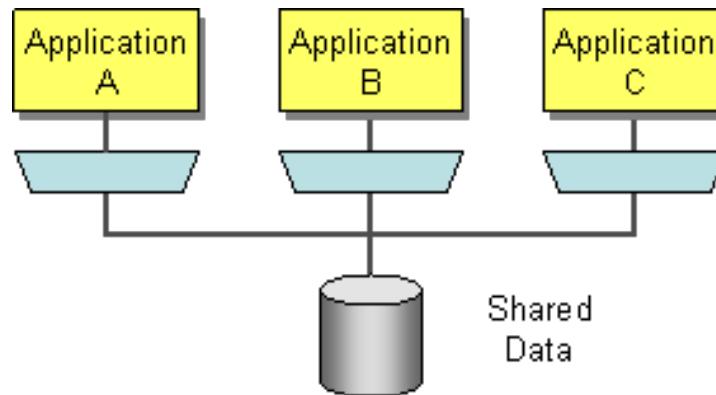




# Polyglot Persistence

# Integration through a shared database

- **Problem:** How can I integrate multiple distributed applications or services so that they work together and can exchange information?



- **Solution:** Integrate applications by having them store their data in a single **Shared Database**.

# Polyglot Persistence

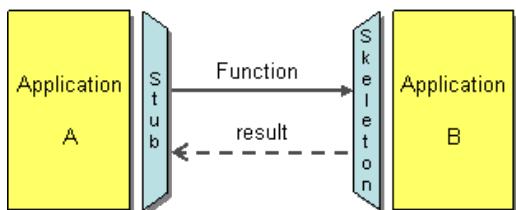


# Integration without a shared Database

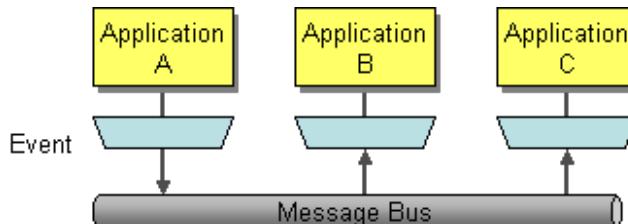


## Service Integration Patterns

### Web API



### Message Queue



### File Transfer Integration

